



**Hewlett Packard**  
Enterprise

# Introduction to perftools

---

LUMI Advanced Workshop

March 5-7, 2025

# Agenda

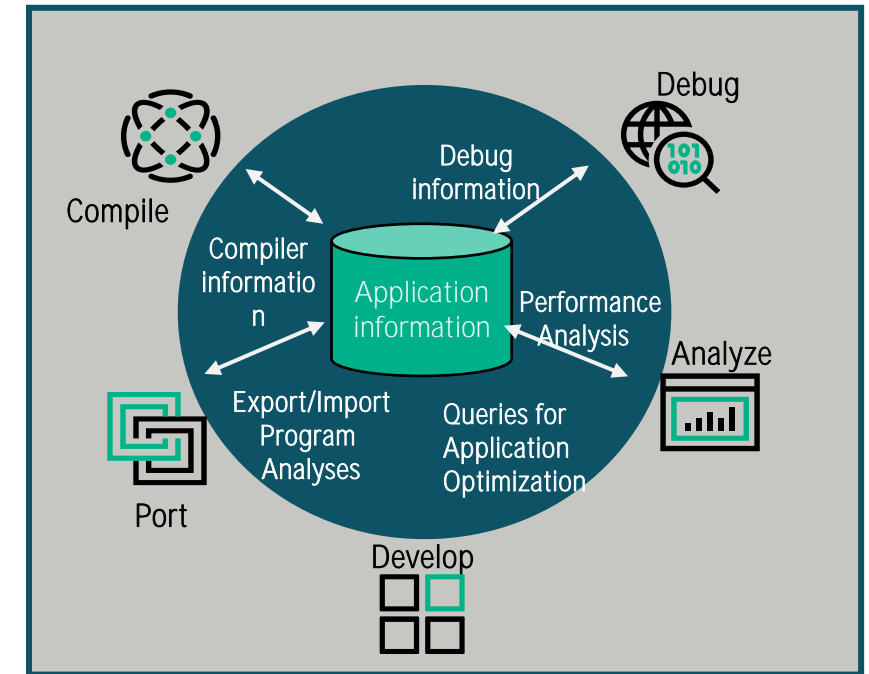
---

- Introduction
  - General Remarks on Performance Analysis
  - Possible Bottlenecks and Remedies
- Perftools
  - Landscape and Properties
  - Sampling vs. Tracing
  - Performance analysis with **perftools-lite**
  - Apprentice2
- Demo



# Motivation

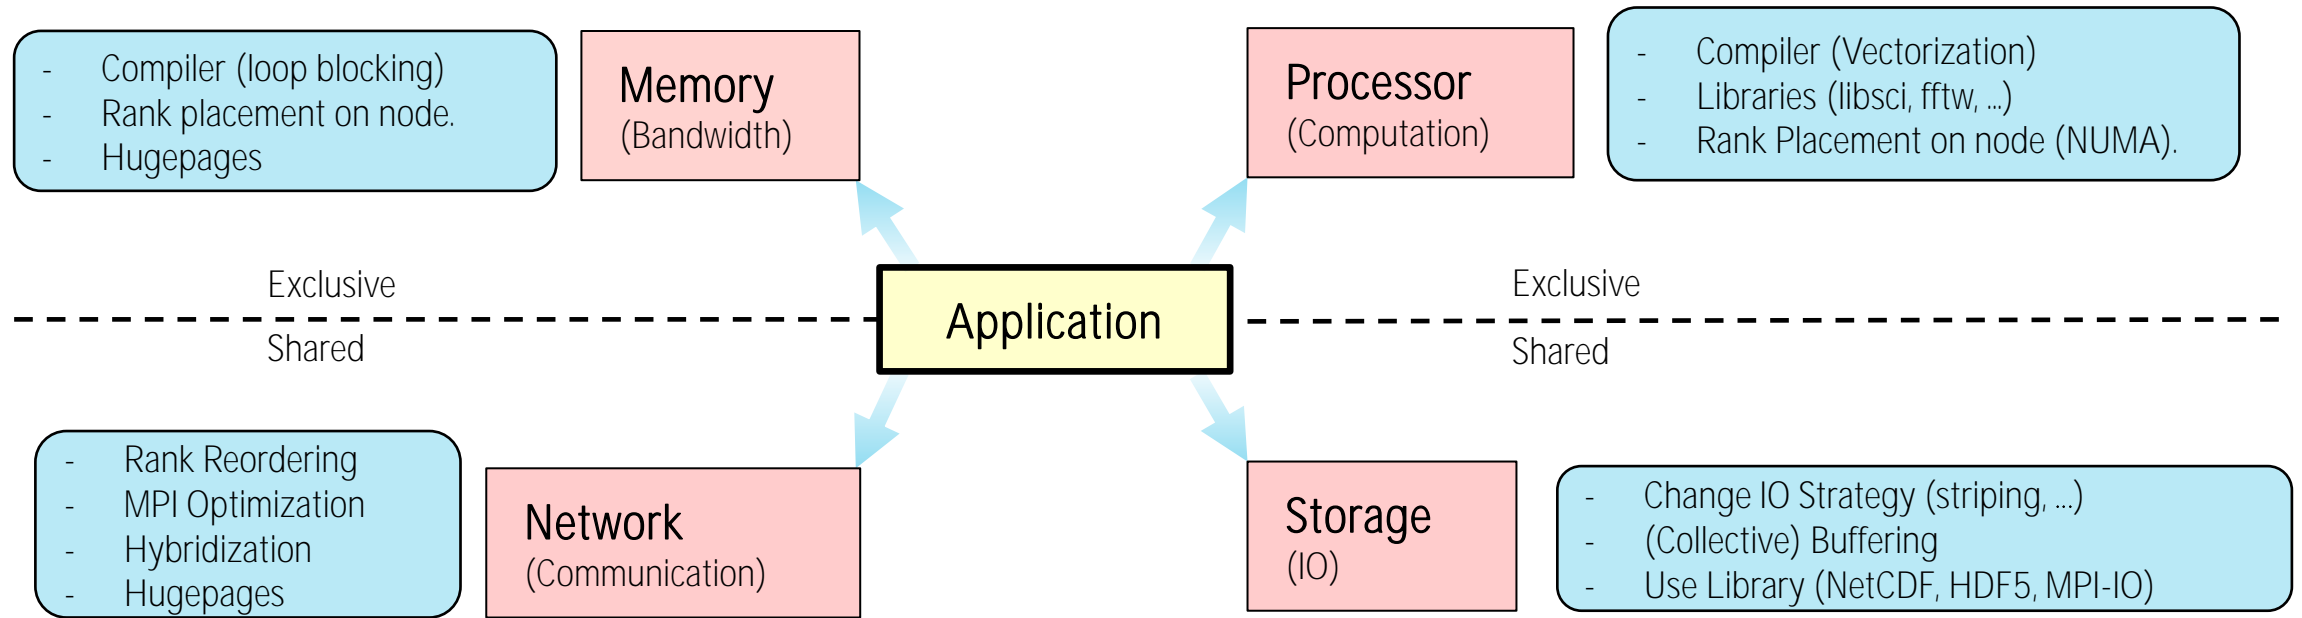
- Very tempting to skip performance analysis when tests validate and time to solution is smaller after port or algorithm improvement.
- But performance does matter!
  - Might still be inefficient most of the of time.
  - Poor code performance can affect other users as well, for example because of bad I / O access patterns.
  - Want to efficiently use expensive resources and get as much information as possible for the allocated resources.
  - Simulating larger models may only be feasible after optimization.
- Applies to various scenarios
  - Code has been ported to a new system or otherwise significantly changed.
  - Application is running in production since a while.
  - Makes extensive use of third-party libraries (distributed ML, dense LA, ...) or even fully proprietary or it is mostly home-grown code.



# General Remarks on Performance Analysis

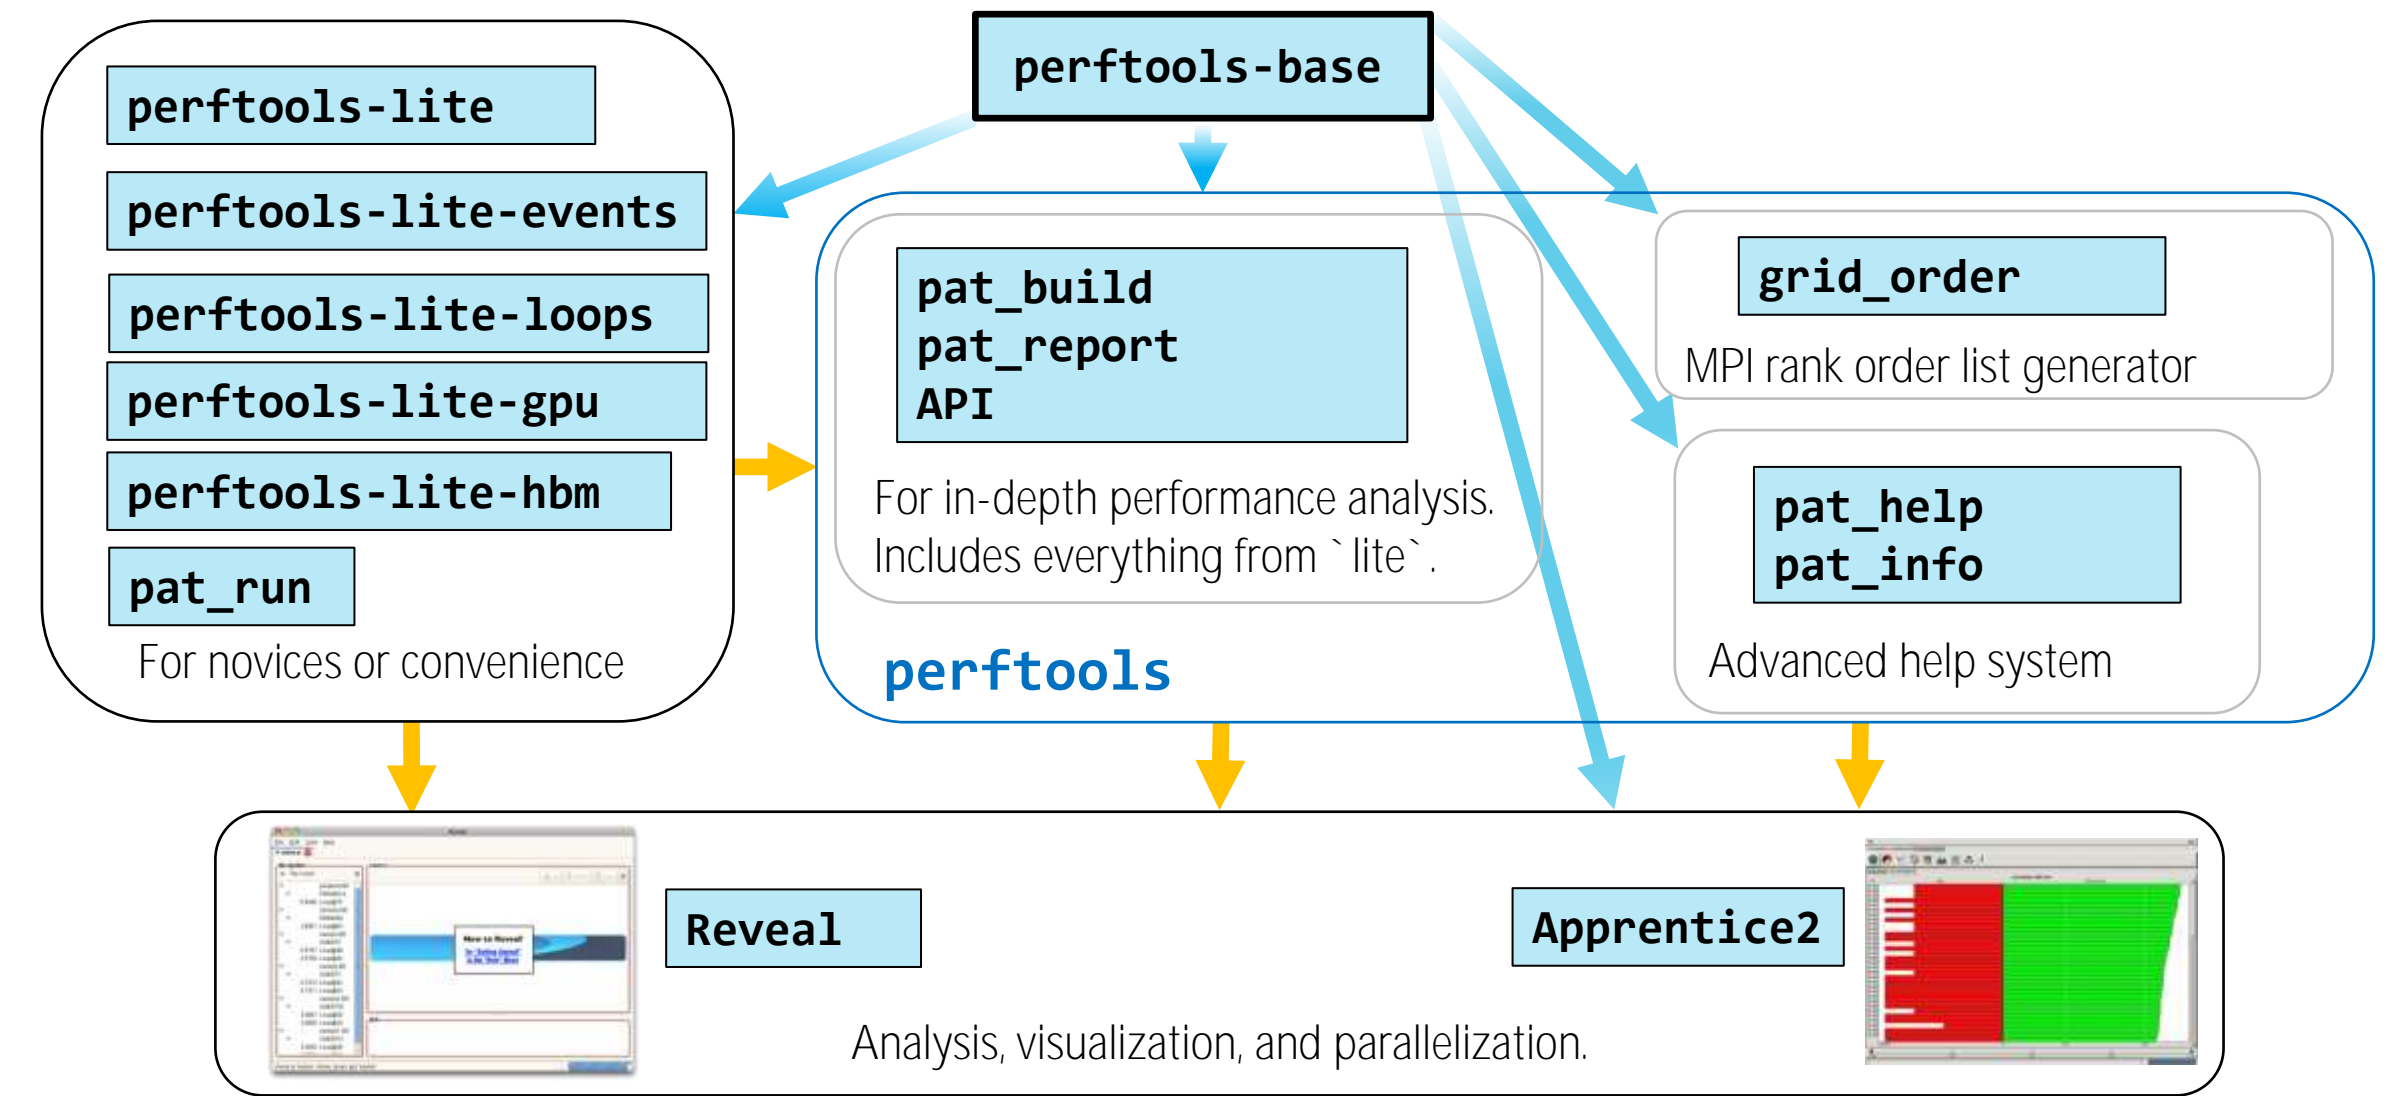
- Performance is usually associated to FLOPs/sec.
  - But what if your code does many scattered memory references like in Graph Analytics and not many FLOPs ?
  - Chose the metric which suits your application (like time to solution or updates/sec instead of FLOPs/sec) and keep that metric throughout the optimization process.
- Use examples with different sizes for your experiments.
  - Small, medium, large, where node count differs by an order of magnitude (strong or weak scaling or both).
    - Try the same model with different grid sizes or number of particles
  - Do not use fully artificial examples but rather meaningful representatives of your target (large scale) simulation.
- 1. Start with low hanging fruits, i.e. avoid code modifications first.
  - Compiler flags, manual rank reordering, optimized libraries, huge pages, use hyperthreads, ...
- 2. Use performance analysis tools
  - Identify critical code regions, guided rank reordering, Automatic parallelization (OpenMP), ...
  - A good understanding of the workflow of your application (Communication, Computation, IO, ... ) helps to better interpret the profiles.

# Bottlenecks and Remedies

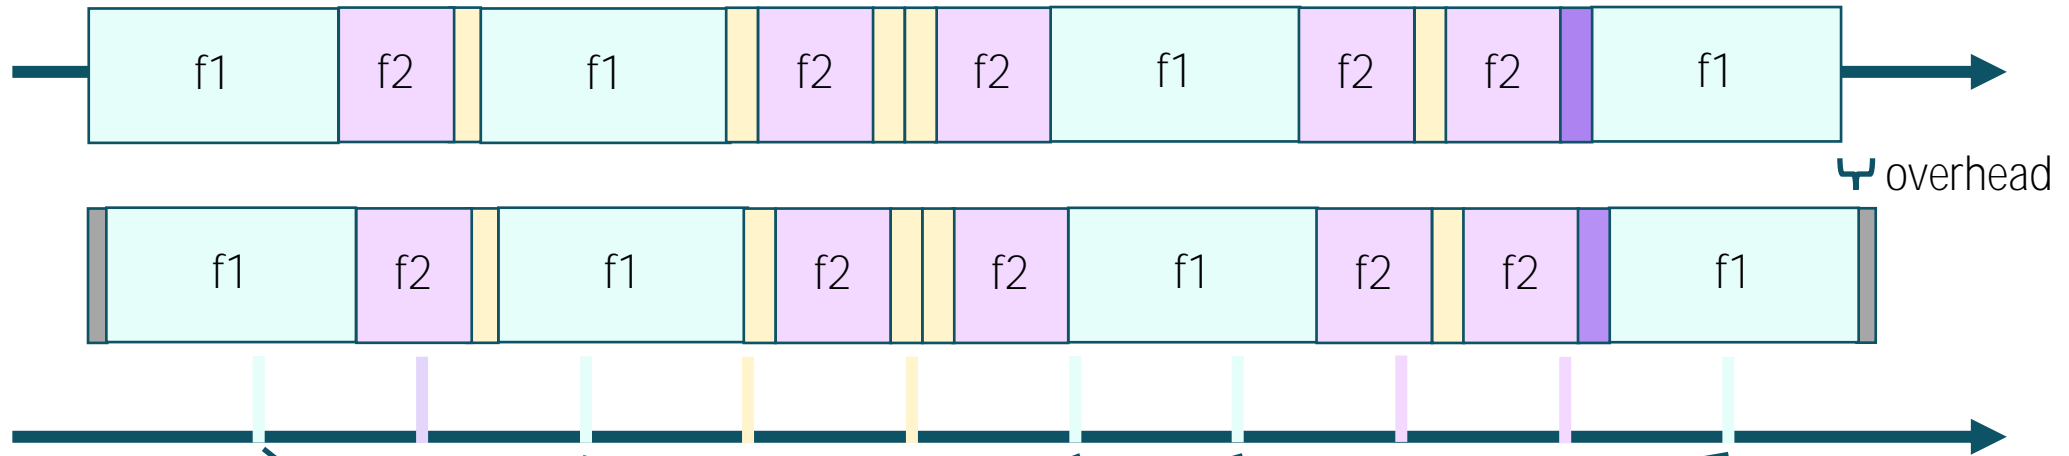


- Good: One bottleneck which can be easily resolved without creating a new one.
- Bad: Several bottlenecks interacting with each other and changing over time.
- Need a profiler to identify bottleneck(s) and a model to estimate optimization potential.

# Perftools Landscape



# First fundamental way of profiling : sampling



Profiling by sampling:

f1	$5/10 = 50\%$
f2	$3/10 = 30\%$
f3	$2/10 = 20\%$
f4	$0/10 = 0\%$ (not displayed)

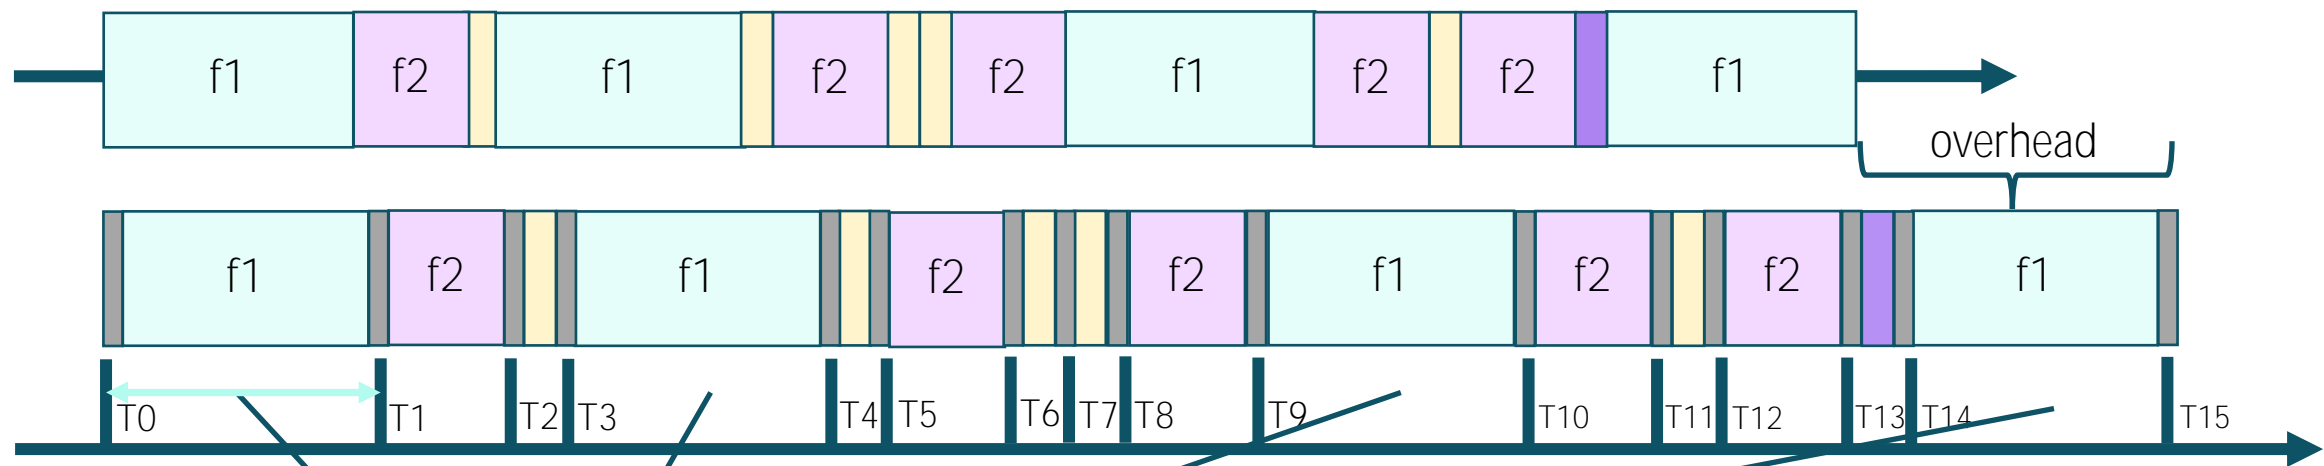
## Advantages

- Only need to instrument main routine
- **Low Overhead** – depends only on sampling frequency
- **Smaller volumes of data produced**

## Disadvantages

- **Only statistical averages available**
- **Some functions may be missed**
- Limited information from performance counters

# Second fundamental way of profiling : tracing



## Profiling by tracing:

f1

$$((T1-T0)+(T4-T3)+(T10-T9)+(T15-T14))) / (T15-T0) \sim 54\% \text{ (50\%)}$$

f2

$$((T2-T1)+ \dots +(T13-T12)) / (T15-T0) \sim 35\% \text{ (30\%)}$$

f3

$$((T3-T2)+ \dots +(T12-T11)) / (T15-T0) \sim 9\% \text{ (20\%)}$$

f4

$$((t14-T13)) / (T15-T0) \sim 2\% \text{ (0\%)}$$

## Advantages

- More accurate and more detailed information
- Data collected from every traced function call

## Disadvantages

- Overheads increases with number of function calls
- Huge volumes of data generated

Sampling outputs (xx%)



# Other Experiments

- Guided Tracing: Combining Sampling and Tracing
  - Trace only functions that are not small (i.e. very few lines of code) and contribute a lot to application's run time.
  - Automatic performance Analysis (APA) is an automated way to do this. (Not intended for [perftools-lite](#))
    1. In a first round identify important routines with a sampling experiment.
    2. Then do a tracing of only these important routines to reduce overhead.
- Loop Work Estimates
  - Special flavor of event tracing targeting loops.
  - Only for Cray programming environment.
  - Useful starting point for porting to multicore or GPUs.
  - Can be used as input to Reveal.
  - Also available with [perftools-lite-loops](#)



# Perftools-lite

---

An easy-to-use version of the Perftools Performance Measurement and Analysis Tool



# Test 1: Generate a Sampling Profile

```
$> module load perftools-base
```

```
$> module load perftools-lite
```

- Subsequent compiler invocations (**cc, CC, ftn**) will automatically insert necessary hooks for profiling (not always up-to-date with latest third-party compilers)

```
$> make clean; make
```

- Resulting executable **app.exe** is automatically instrumented. Object files needed!
- Not instrumented application stored as **app.exe+orig**

```
$> srun -n 8 app.exe >& job.out
```

- The report is printed to stdout.
- Successful execution creates a **app.exe+\*/** directory for further analysis. Change this directory name with the **PAT\_RT\_EXPDIR\_NAME** environment variable.
- Use the environment variable **PAT\_RT\_SAMPLING\_INTERVAL** to change sampling interval (default is 10000 microseconds)

# Output: Sampling

## Regular program output

```
CrayPat/X: Version 22.06.0 Revision 4b5ab6256 05/21/22 02:03:49
Sequential version array size
  mimax = 513 mjmax = 513 mkmax = 1025
Parallel version array size
  mimax = 259 mjmax = 259 mkmax = 515
  imax = 257 jmax = 257 kmax = 513
I-decomp = 2 J-decomp = 2 K-decomp = 2
Start rehearsal measurement process.
Measure the performance in 100 times.

MFLOPS: 15638.324504 time(s): 5.779363 2.794488e-04

Now, start the actual measurement process.
The loop will be excuted in 50 times
This will take about one minute.
Wait for a while

cpu : 28.858915 sec.
Loop executed for 50 times
Gosa : 2.763023e-04
MFLOPS measured : 15658.861128
Score based on Pentium III 600MHz : 189.025364
```

```
#####
#                               #
#   CrayPat-lite Performance Statistics   #
#                               #
#####

CrayPat/X: Version 22.06.0 Revision 4b5ab6256 05/21/22 02:03:49
Experiment:      lite lite/sample_profile
Number of PEs (MPI ranks):      8
Numbers of PEs per Node:      8
Numbers of Threads per PE:      1
Number of Cores per Socket:    16

Execution start time: Sat Aug 20 23:43:46 2022
System name and speed: nid005017 2.626 GHz (nominal)
AMD Trento      CPU Family: 25 Model: 48 Stepping: 1
Core Performance Boost: All 8 PEs have CPB capability

Avg Process Time:      36.92 secs
High Memory:      15,544.6 MiBytes      1,943.1 MiBytes per PE
I/O Write Rate:      10.359011 MiBytes/sec
```

## General job information

## Sampling summary

Notes for table 1:

This table shows functions that have significant exclusive sample hits, averaged across ranks.  
For further explanation, use: pat\_report -v -O samp\_profile ...

Table 1: Profile by Function

Samp%	Samp	Imb. Samp	Imb. Samp%	Group
				Function=[MAX10] PE=HIDE
100.0%	3,693.2	--	--	Total
84.1%	3,107.8	--	--	USER
79.6%	2,938.8	65.2	2.5%	jacobi
4.6%	169.0	2.0	1.3%	initmt
14.7%	541.5	--	--	ETC
13.3%	490.2	31.8	7.0%	__cray_memcpy_ROME
1.4%	50.9	4.1	8.6%	__cray_memset_ROME
1.1%	41.8	--	--	MPI

...

## Further analysis

Notes for table 3:

This table shows energy and power usage for the nodes with the maximum, mean, and minimum usage, as well as the sum of usage over all nodes.  
Energy and power for accelerators is also shown, if applicable.  
For further explanation, use: pat\_report -v -O program\_energy ...

Table 3: Program energy and power usage (from Cray PM)

Node Energy (J)	Node Power (W)	Process Time	PE=HIDE
25,504	687.169	37.114590	Total

Notes for table 4:

This table show the average time and number of bytes written to each output file, taking the average over the number of ranks that wrote to the file. It also shows the number of write operations, and average rates.  
For further explanation, use: pat\_report -v -O write\_stats ...

Table 4: File Output Stats by Filename

Avg	Avg	Write Rate	Number	Avg	Bytes/	File
Name=!	Write	Write	of	Writes	Call	PE=HIDE
Time per	MiBytes	MiBytes/sec	Writer	per		
Writer	per		Ranks	Writer		
Rank	Rank			Rank		
0.000052	0.000582	11.224027	1	18.0	33.89	stdout

...

===== End of CrayPat-lite output =====

## Test 2: Generate an Event Profile

```
$> module sw perftools-lite perftools-lite-events
```

- If **perftools-lite** module not loaded, load subsequently **perftools-base** and **perftools-lite-events**.

```
$> rm app.exe; make
```

- Only relink of **app.exe** necessary if object files and user libraries have been generated with another **perftools-lite\*** module.
- Otherwise do a **make clean; make**

```
$> srun -n 8 app.exe >& job.out
```

- The report is printed to stdout.
- Successful execution creates a **app.exe+\*/** directory for further analysis.

# Output: Event Tracing

## Regular program output

```
CrayPat/X: Version 22.06.0 Revision 4b5ab6256
05/21/22 02:03:49
Sequential version array size
  mimax= 1025  mjmax= 513  mkmax= 513
Parallel version array size
  mimax= 515  mjmax= 259  mkmax= 259
  imax= 513  jmax= 257  kmax= 257
  I-decomp= 2  J-decomp= 2  K-decomp= 2

Start rehearsal measurement process.
Measure the performance in 3 times.
  MFLOPS: 15604.522995096944  time(s):
5.7918816709999987, 4.324619949E-4
Now, start the actual measurement process.
The loop will be excuted in 50 times.
This will take about one minute.
Wait for a while.
Loop executed for 50 times
Gosa : 4.21404955E-4
MFLOPS: 15611.16692542893  time(s): 28.947083569
Score based on Pentium III 600MHz : 188.449631
```

```
#####
#
#   CrayPat-lite Performance Statistics   #
#
#####
CrayPat/X: Version 22.06.0 Revision 4b5ab6256 05/21/22
02:03:49
Experiment:                lite  lite-events
Number of PEs (MPI ranks):      8
Numbers of PEs per Node:       8
Numbers of Threads per PE:      1
Number of Cores per Socket:     64
Execution start time:  Fri Nov 18 16:02:17 2022
System name and speed:  nid001050  2.078 GHz (nominal)
AMD  Milan                CPU Family: 25  Model: 1
Stepping: 1
Core Performance Boost: All 8 PEs have CPB capability

Avg Process Time:    37.05 secs
High Memory:         15,667.4 MiBytes    1,958.4 MiBytes per PE
I/O Write Rate:      3.182425 MiBytes/sec
```

## General job information

## Event tracing summary. Note difference to sampling

Notes for table 1:

This table shows functions that have significant exclusive time, averaged across ranks.  
For further explanation, use: `pat_report -v -O profile ...`

Table 1: Profile by Function Group and Function

Time%	Time	Imb. Time	Imb. Time%	Calls	Group Function=[MAX10] PE=HIDE
100.0%	37.007250	--	--	1,461.4	Total
98.4%	36.410057	--	--	3.0	USER
76.9%	28.457104	0.371909	1.5%	1.0	jacobi
15.4%	5.690022	0.078645	1.6%	1.0	himenobmtxp
6.1%	2.262931	0.004594	0.2%	1.0	initmt
1.6%	0.585681	--	--	983.0	MPI
1.3%	0.482364	0.614413	64.0%	180.0	MPI_WAITALL

Observation: MPI utilization

No suggestions were made because all ranks are on one node.

Notes for table 2:

This table shows energy and power usage for the nodes with the maximum, mean, and minimum usage, as well as the sum of usage over all nodes.  
Energy and power for accelerators is also shown, if applicable.  
For further explanation, use: `pat_report -v -O program_energy ...`

Table 2: Program energy and power usage (from Cray PM)

Node Energy (J)	Node Power (W)	Process Time	PE=HIDE
9,055	244.407	37.048880	Total

Notes for table 3:

This table show the average time and number of bytes written to each output file, taking the average over the number of ranks that wrote to the file. It also shows the number of write operations, and average rates.

For further explanation, use: `pat_report -v -O write_stats ...`

Table 3: File Output Stats by Filename

Avg Write Time per Writer Rank	Avg Write MiBytes per Writer Rank	Write Rate MiBytes/sec	Number of Writer Ranks	Avg Writes per Writer Rank	Avg Bytes/Call	File
0.000024	0.000008	0.321773	8	1.0	8.00	stderr
0.000022	0.000613	27.644603	1	18.0	35.72	stdout

Program invocation: himeno.exe

For a complete report with expanded tables and notes, run:  
`pat_report /pfs/lustrep2/projappl/project_465000297/alfiolaz/work/perftools-lite/test/expfile.lite-events.2046565`

For help identifying callers of particular functions:  
`pat_report -O callers+src /pfs/lustrep2/projappl/project_465000297/alfiolaz/work/perftools-lite/test/expfile.lite-events.2046565`  
To see the entire call tree:  
`pat_report -O calltree+src /pfs/lustrep2/projappl/project_465000297/alfiolaz/work/perftools-lite/test/expfile.lite-events.2046565`

For interactive, graphical performance analysis, run:  
`app2 /pfs/lustrep2/projappl/project_465000297/alfiolaz/work/perftools-lite/test/expfile.lite-events.2046565`

===== End of CrayPat-lite output =====

## Further analysis

## Test 3: Generate a Loop Profile (CCE only)

```
$> module sw perftools-lite perftools-lite-loops
```

- If perftools-lite module not loaded, load subsequently perftools-base and perftools-lite-loops. Only for **PrgEnv-cray**.

```
$> make clean; make
```

- Need to clean everything and rebuild.
- Compiler drivers will use **-h profile\_generate** for Fortran or **-finstrument-loops** for C implicitly. This flag turns off OpenMP and significant compiler loop restructuring optimizations except for vectorization.

```
$> srun -n 8 app.exe >& job.out
```

- Successful execution creates a **app.exe+\*/** directory for further analysis.
- The report is printed to stdout.

# Output: Loop Profile

## Regular program output

```
CrayPat/X: Version 22.06.0 Revision 4b5ab6256 05/21/22 02:03:49
Sequential version array size
mimax = 513 mjmax = 513 mkmax = 1025
Parallel version array size
mimax = 259 mjmax = 259 mkmax = 515
imax = 257 jmax = 257 kmax =513
I-decomp = 2 J-decomp = 2 K-decomp =2
Start rehearsal measurement process.
Measure the performance in 100 times.

MFLOPS: 15652.801400 time(s): 5.774017 2.794488e-04

Now, start the actual measurement process.
The loop will be excuted in 50 times
This will take about one minute.
Wait for a while

cpu : 28.869171 sec.
Loop executed for 50 times
Gosa : 2.763023e-04
MFLOPS measured : 15653.298226
Score based on Pentium III 600MHz : 188.958211
```

```
#####
#
#          CrayPat-lite Performance Statistics          #
#
#####

CrayPat/X: Version 22.06.0 Revision 4b5ab6256 05/21/22 02:03:49
Experiment:      lite lite-loops
Number of PEs (MPI ranks):      8
Numbers of PEs per Node:      8
Numbers of Threads per PE:      1
Number of Cores per Socket:      64
Accelerator Model: AMD MI200 Memory: 32.00 GB Frequency: 1.09 GHz

Execution start time: Sun Aug 21 00:29:19 2022
System name and speed: nid005013 2.361 GHz (nominal)
AMD Trento          CPU Family: 25 Model: 48 Stepping: 1
Core Performance Boost: All 8 PEs have CPB capability

Avg Process Time:      37.12 secs
High Memory:      15,593.3 MiBytes 1,949.2 MiBytes per PE
```

## General job information

## Loop Statistics by function

Notes for table 1:

This table shows a nested view of loops that have significant inclusive time, averaged across ranks. Intervening function calls are not shown (as if all functions were inlined). For each loop, the table shows its inclusive time, the number of times it was executed, and the average number of iterations for each execution. Times in this table include overhead from loop instrumentation. For an alternative view that shows min and max iterations, and exclusive times, use the option: `-O loop_times`. For further explanation, use: `pat_report -v -O loop_nest ...`

Table 1: Nested view of Loop Inclusive Time

	Incl Time%	Incl Time	Loop Exec	Loop Trips Avg	Calltree=/[.]LOOP[.] PE=HIDE
	100.0%	37.08	--	--	Total
1	93.4%	34.64	2	30.0	jacobi.LOOP.1.li.236
2	79.0%	29.31	60	255.0	jacobi.LOOP.2.li.240
3	79.0%	29.31	15,300	255.0	jacobi.LOOP.3.li.241
4	78.2%	29.00	3,901,500	511.0	jacobi.LOOP.4.li.242
2	13.5%	5.02	60	255.0	jacobi.LOOP.5.li.263
3	13.5%	5.02	15,300	255.0	jacobi.LOOP.6.li.264
4	9.6%	3.54	3,901,500	511.0	jacobi.LOOP.7.li.265
1	4.3%	1.61	1	259.0	initmt.LOOP.4.li.191
2	4.3%	1.61	259	259.0	initmt.LOOP.5.li.192
3	4.3%	1.59	67,081	515.0	initmt.LOOP.6.li.193
1	2.2%	0.82	1	257.0	initmt.LOOP.1.li.210
2	2.2%	0.82	257	257.0	initmt.LOOP.2.li.211
3	2.1%	0.79	66,049	513.0	initmt.LOOP.3.li.212

Program invocation: himeno.exe

For a complete report with expanded tables and notes, run:  
`pat_report /pfs/lustrep3/users/lazzaroa/exercises/perftools-lite/C/himeno.exe+44287-8745533t`

For help identifying callers of particular functions:  
`pat_report -O callers+src /pfs/lustrep3/users/lazzaroa/exercises/perftools-lite/C/himeno.exe+44287-8745533t`  
To see the entire call tree:  
`pat_report -O calltree+src /pfs/lustrep3/users/lazzaroa/exercises/perftools-lite/C/himeno.exe+44287-8745533t`

For interactive, graphical performance analysis, run:  
`app2 /pfs/lustrep3/users/lazzaroa/exercises/perftools-lite/C/himeno.exe+44287-8745533t`

===== End of CrayPat-lite output =====

Subroutine

Line number

Nested Loops



# Generate an Event Profile for GPU Experiments

```
$> module load perftools-lite-gpu
```

- If **perftools-lite** module not loaded, load subsequently **perftools-base** and **perftools-lite-gpu**.

```
$> rm app.exe; make
```

- Only relink of **app.exe** necessary if object files and user libraries have been generated with another **perftools-lite\*** module.
- Otherwise do a **make clean; make**

```
$> srun -n 8 app.exe >& job.out
```

- The report is printed to stdout.
- Successful execution creates a **app.exe+\*/** directory for further analysis.

# Output: GPU Profile (MPI+ompenMP offload)

## Regular program output

```
CrayPat/X: Version 24.03.0 Revision 1705b3005 sles15.5_x86_64
02/20/24 20:53:05
Sequential version array size
  mimax = 513 mjmax = 513 mkmax = 1025
Parallel version array size
  mimax = 259 mjmax = 259 mkmax = 515
  imax = 257 jmax = 257 kmax =513
I-decomp = 2 J-decomp = 2 K-decomp =2
Start rehearsal measurement process.
Measure the performance in 10 times.

MFLOPS: 8127.332464 time(s): 11.120444 4.351904e-04

Now, start the actual measurement process.
The loop will be excuted in 50 times
This will take about one minute.
Wait for a while

cpu : 54.424242 sec.
Loop executed for 50 times
Gosa : 4.239466e-04
MFLOPS measured : 8303.243579
Score based on Pentium III 600MHz : 100.232298
```

```
#####
#          CrayPat-lite Performance Statistics          #
#          #####                                         #
CrayPat/X: Version 24.03.0 Revision 1705b3005 sles15.5_x86_64 02/20/24
20:53:05
Experiment:      lite  lite-gpu
Number of PEs (MPI ranks):      8
Numbers of PEs per Node:      8
Numbers of Threads per PE:      1
Number of Cores per Socket:      64
Accelerator Model: AMD MI200 Series; Memory: 64.00 GB; Frequency: 1.70 GHz

Execution start time: Wed Jan  8 11:09:12 2025
System name and speed: nid005976 2.000 GHz (nominal)
AMD Trento CPU Family: 25 Model: 48 Stepping: 1
Core Performance Boost: All 8 PEs have CPB capability

Avg Process Time:      67.78 secs
High Memory:      15,162.1 MiBytes      1,895.3 MiBytes per PE
I/O Read Rate:      9,722530 MiBytes/sec
I/O Write Rate:      7.356752 MiBytes/sec
```

## General job information

## gpu profiling summary.

Table 1: Caller View by Acc Id of Accelerator Time and Bytes Transferred

Time%	Time	Function Caller PE=HIDE
100.0%	67.624946	Total
88.4%	59.774645	__omp_offloading_73ac72ce_d201d046_jacobi_1242_debug___.REGION@li.244
		__omp_offloading_73ac72ce_d201d046_jacobi_1242_debug___.REGION@li.244(ovhd):himeno.c:line.244
7.9%	5.361571	__omp_offloading_73ac72ce_d201d046_jacobi_1268_debug___.REGION@li.270
		__omp_offloading_73ac72ce_d201d046_jacobi_1268_debug___.REGION@li.270(ovhd):himeno.c:line.270
3.0%	2.055567	initmt
		main:himeno.c:line.102

Table 5: File Output Stats by Filename

Avg Write Time per Writer Rank	Avg Write MiBytes per Writer Rank	Write Rate MiBytes/sec	Number of Writer Ranks	Avg Writes per Writer Rank	Bytes/ Call	File Name=!\x/^/(proc sys)/ PE=HIDE
0.000080	0.000587	7.356752	1	18.0	34.17	stdout

Table 3: Program Energy and Power Usage from Cray PM

PE=HIDE

=====			
Total			
-----			
PM Energy Node	563 W	38,137 J	
PM Energy Cpu	84 W	5,672 J	
PM Energy Memory	78 W	5,287 J	
PM Energy Acc0	98 W	6,630 J	
PM Energy Acc1	92 W	6,256 J	
PM Energy Acc2	95 W	6,452 J	
PM Energy Acc3	93 W	6,300 J	
Process Time		67.775526 secs	
=====			

Program invocation: himeno.exe

For a complete report with expanded tables and notes, run:

```
pat_report
/pfs/lustrep4/projappl/project_462000031/bracconi/Training/training_exercises/p
erftools-lite-gpu/C/expfile..9012372
```

For help identifying callers of particular functions:

```
pat_report -O callers+src
/pfs/lustrep4/projappl/project_462000031/bracconi/Training/training_exercises/p
erftools-lite-gpu/C/expfile..9012372
```

To see the entire call tree:

```
pat_report -O calltree+src
/pfs/lustrep4/projappl/project_462000031/bracconi/Training/training_exercises/p
erftools-lite-gpu/C/expfile..9012372
```

For interactive, graphical performance analysis, run:

```
app2
/pfs/lustrep4/projappl/project_462000031/bracconi/Training/training_exercises/p
erftools-lite-gpu/C/expfile..9012372
```

## Further analysis

# Observations and Remarks

- No intervention needed for build system and batch scripts.
  - Only make sure to use the compiler driver wrappers **CC**, **cc**, and **ftn**.
- What we did not see with these simple tests:
  - **perftools-lite** can produce rank reordering files for MPI to optimize the communication. Not visible here because of small portion of time spent in communication or job size.
  - The resulting **app+exe\*/** directory can be processed with **pat\_report**, Apprentice2, and Reveal for further analysis. From the sample experiment one can retrieve hardware performance counter information.
- Tailored profiling, i.e. for specific routines, trace groups, or specific portions of the code is not possible.
  - Need the regular **perftools** module for this in-depth analysis.
- **CRAYPAT\_LITE** environment variable can be used to distinct output files.
- Record Subset of PEs during execution: **export PAT\_RT\_EXPFILE\_PES=0,4,5,10**
- Use **CRAYPAT\_LITE\_WHITELIST** for binaries you DO want instrumented (rest ignored).



## Further analysis (without re-running)

- Generate full report

```
$> pat_report app.exe+pat*/ > rpt
```

- Generate report with call tree (or by callers)

```
$> pat_report -O calltree+src
```

```
$> pat_report -O callers+src
```

- Show each MPI rank or each OpenMP thread in report

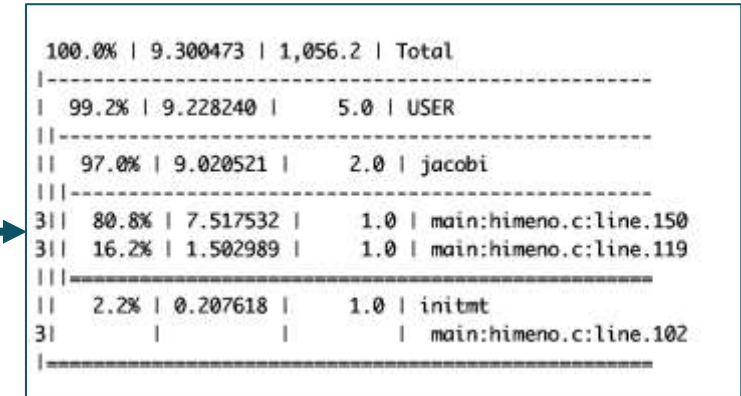
```
$> pat_report -s pe=ALL
```

```
$> pat_report -s th=ALL
```

- Generate a preview of data before processing the full report

```
$> pat_report -Q1
```

- Produces report from single (lexically first) '.ap2' file
- Useful for jobs with large number of processes



100.0%	9.300473	1,056.2	Total
99.2%	9.228240	5.0	USER
97.0%	9.020521	2.0	jacobi
80.8%	7.517532	1.0	main:himeno.c:line.150
16.2%	1.502989	1.0	main:himeno.c:line.119
2.2%	0.207618	1.0	initmt
			main:himeno.c:line.102

# Further analysis (without re-running)

- Generate report from specific subset of ranks

```
$> pat_report -s filter_input='pe==0'
```

- Report with only PE 0 data

```
$> pat_report -s filter_input='pe<5'
```

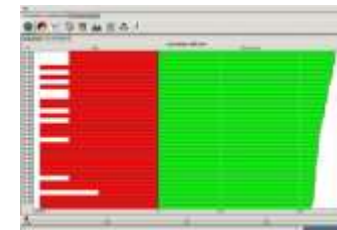
- Report with data from first 5 ranks
- Use `pat_help report filtering` for more details

- Don't see an expected function?

- Use the `pat_report -P` option to disable pruning.
- You should be able to see the caller/callee relationship with `pat_report -P -O callers`
- Use `'pat_report -T'` to see functions that didn't take much time
- Still don't see it? Check the compiler listing to see if the function was inlined.

- Also try the GUI for analyzing performance analysis results.

```
$> app2 app.exe+*/
```



# On-node analysis

- Memory bandwidth sensitivity guidance

Functions Slowed By Memory Bandwidth Utilization

The performance data for the functions shown below suggest that their performance is limited by memory bandwidth. To confirm this, try running with fewer processes placed on each node.

Samp%	Memory	Stall	Function
	Traffic	PerCent	Numanode=HIDE
	/		PE=HIDE
	Nominal		
	Peak		
40.9%	54.1%	93.8%	daxpy_kernel_8
36.1%	59.4%	93.8%	dgemv_kernel_4x4



# On-node analysis

- Example traffic from an MPI+OpenMP run.

Table 3: Memory Bandwidth by Numanode (limited entries shown)

Available in default report assuming processors supports collecting the data

Memory Traffic GBytes	Local Memory Traffic GBytes	Remote Memory Traffic GBytes	Thread Time	Memory Traffic GBytes / Sec	Memory Traffic / Nominal Peak	Numanode Node Id=[max3,min3] PE=HIDE Thread=HIDE
184.47	173.59	10.89	11.578777	15.93	20.7%	numanode.0
183.50	173.59	9.91	11.569322	15.86	20.7%	nid.63
182.61	172.40	10.21	11.578777	15.77	20.5%	nid.61
178.55	167.75	10.80	11.563156	15.44	20.1%	nid.71
178.10	168.14	9.96	11.562097	15.40	20.1%	nid.62
178.08	168.07	10.01	11.564512	15.40	20.1%	nid.68
178.01	167.20	10.82	11.572032	15.38	20.0%	nid.70
60.36	14.73	45.62	9.073119	6.65	8.7%	numanode.1
60.36	14.73	45.62	9.072693	6.65	8.7%	nid.63
59.88	14.33	45.55	9.071553	6.60	8.6%	nid.62
59.48	14.19	45.29	9.068044	6.56	8.5%	nid.68
58.78	13.70	45.08	9.069259	6.48	8.4%	nid.70
58.67	13.87	44.81	9.071591	6.47	8.4%	nid.69
58.53	13.86	44.67	9.067146	6.46	8.4%	nid.71

Notice remote memory traffic by OpenMP threads



# On-node analysis

- Low vectorization guidance.

Functions with Low Vectorization

The performance data for the functions shown below suggest that their performance could be improved by increased vectorization. Use compiler optimization messages to identify loops in those functions that were not vectorized and try to use directives or restructure the loops to enable them to vectorize.

Samp%	Vector	Stall	Function
	intensity	PerCent	PE=HIDE
			Thread=HIDE
-----			
47.7%	0.3%	15.2%	depose_jxjyz_esirkepov_1_1_1_
=====			





# On-node analysis

- Memory latency sensitivity guidance.

Functions Slowed By Memory Latency

The performance data for the functions shown below suggest that their performance is limited by memory latency. It could be beneficial to modify prefetching in loops in those functions, by modifying compiler-generated prefetches or inserting directives into the source code.

Samp%	Memory	Stall	Function
	Traffic	PerCent	Numanode=HIDE
	/		PE=HIDE
	Nominal		
	Peak		
-----			
72.8%	34.8%	33.9%	dim3_sweep\$dim3_sweep_module_
=====			



# Documentation

- Module help
  - module help perftools-base
  - module help perftools
  - module help perftools-lite
  - module help perftools-lite-\*
- Man pages
  - man pat\_build
  - man pat\_report
- Advanced help system
  - pat\_help
  - pat\_info [[.ap2\_file] [experiment\_data\_directory]...]
- Web pages:
  - <https://cpe.ext.hpe.com/docs/24.03/performance-tools/index.html>



# APPRENTICE2

---

Display your performance analysis results

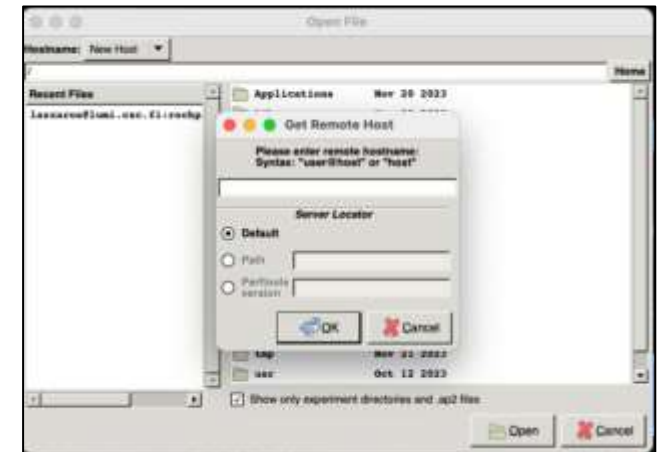


# Apprentice2 – Desktop installer with remote access to LUMI

- Remote visualization of the PAT data
- Installers for Windows and MacOS available at
  - [/opt/cray/pe/perftools/<version>/share/desktop\\_installers/](/opt/cray/pe/perftools/<version>/share/desktop_installers/)
  - No Linux package!
  - Note that the Windows version doesn't support ssh key authentication
    - Cannot be used for LUMI
- MacOS application startup window:



- Open PAT data on LUMI via **File -> Open Remote...**
  - Insert the **<login>@lumi.csc.fi** with **Default** path
  - Navigate to your PAT data directories
  - Click to open the directory of the PAT experiment



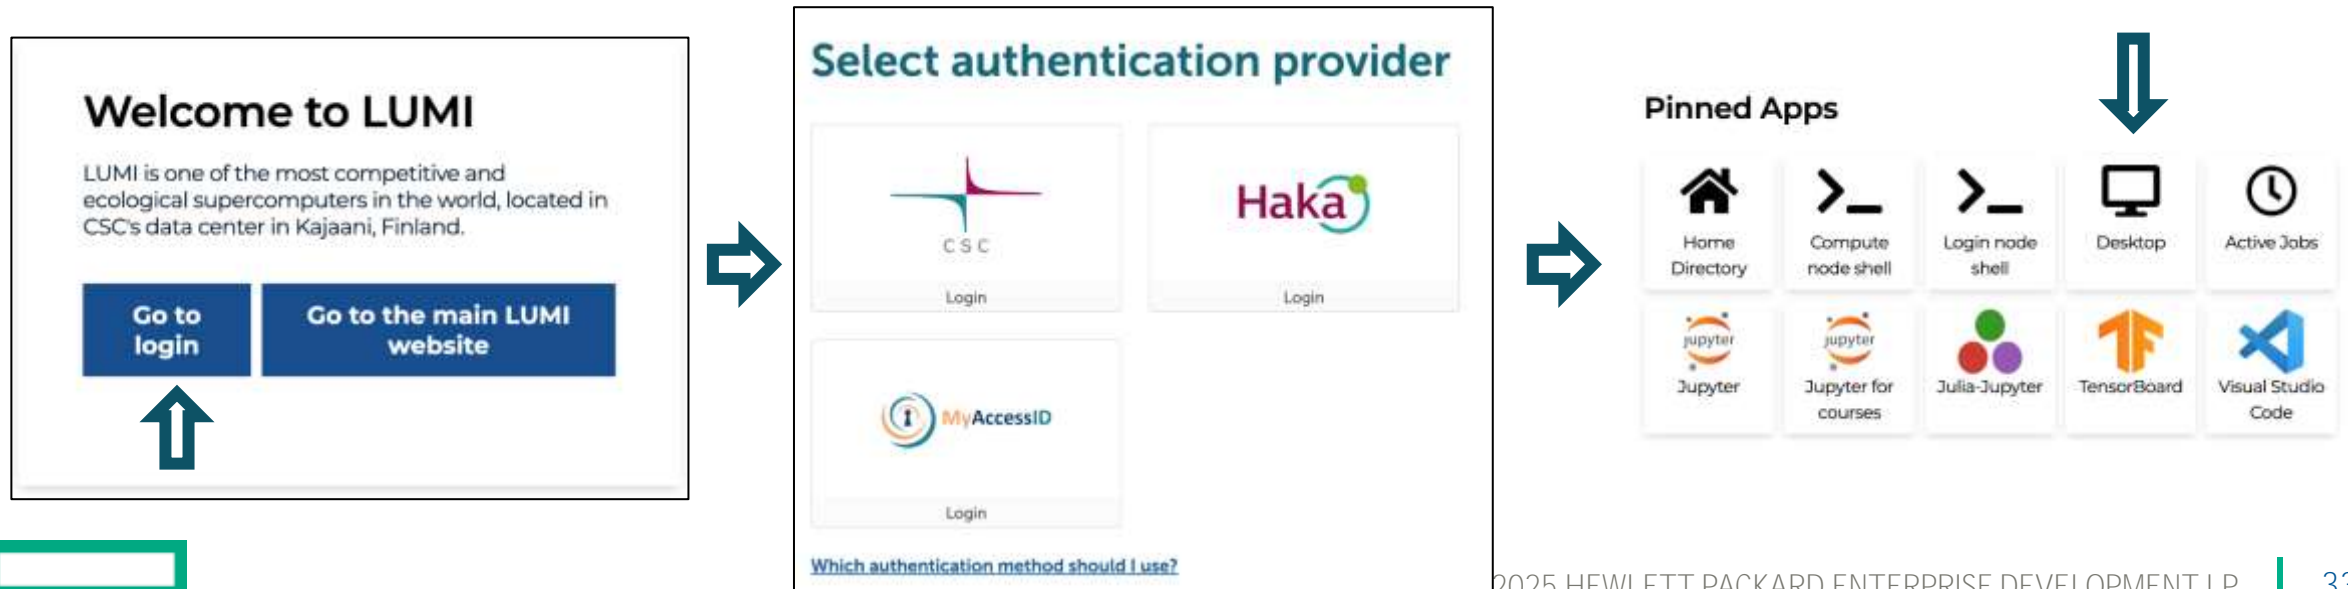
# Apprentice2 – Desktop installer with copy of the PAT data

- It is possible to copy the PAT directories on your laptop/desktop
  - Depending on the experiment, it can be quite large
- It works on Windows



# Apprentice2 – LUMI GUI export (1)

- X export from LUMI (e.g. `ssh -X`) possible but it can be slow
- A better solution is to use VNC from LUMI
  - More info via **module spider VNC**
- LUMI provides a new user interface (suggested solution)
  - <https://www.lumi-supercomputer.eu/introducing-a-new-web-interface-for-lumi/>
  - It offers Open OnDemand as a browser-based interface to launch interactive applications at <https://www.lumi.csc.fi/public/>
    - Follow the instructions at <https://docs.lumi-supercomputer.eu/runjobs/webui/>

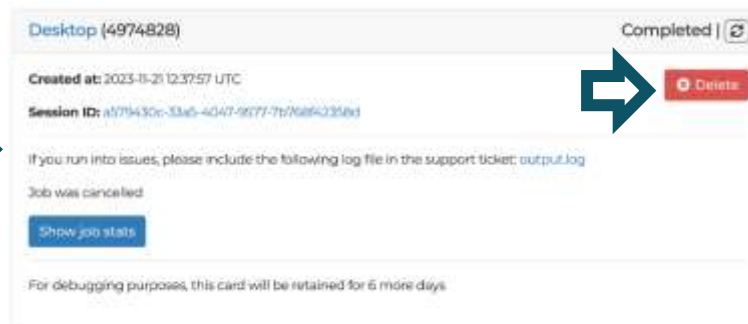
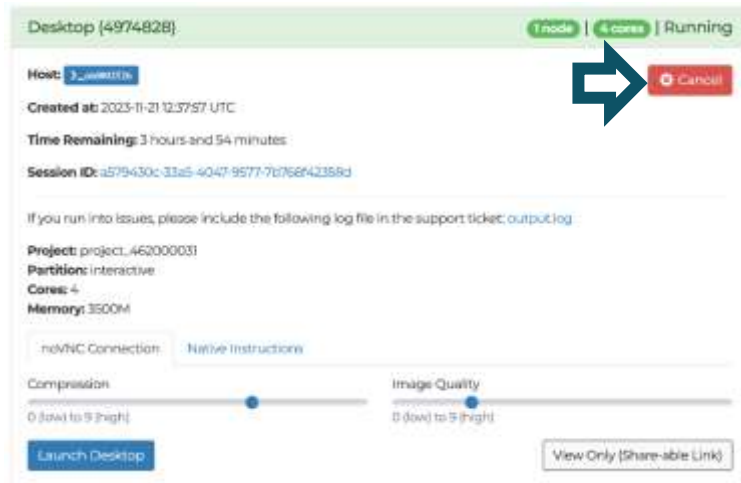


# Apprentice2 – LUMI GUI export (2)



# Apprentice2 – LUMI GUI export (3)

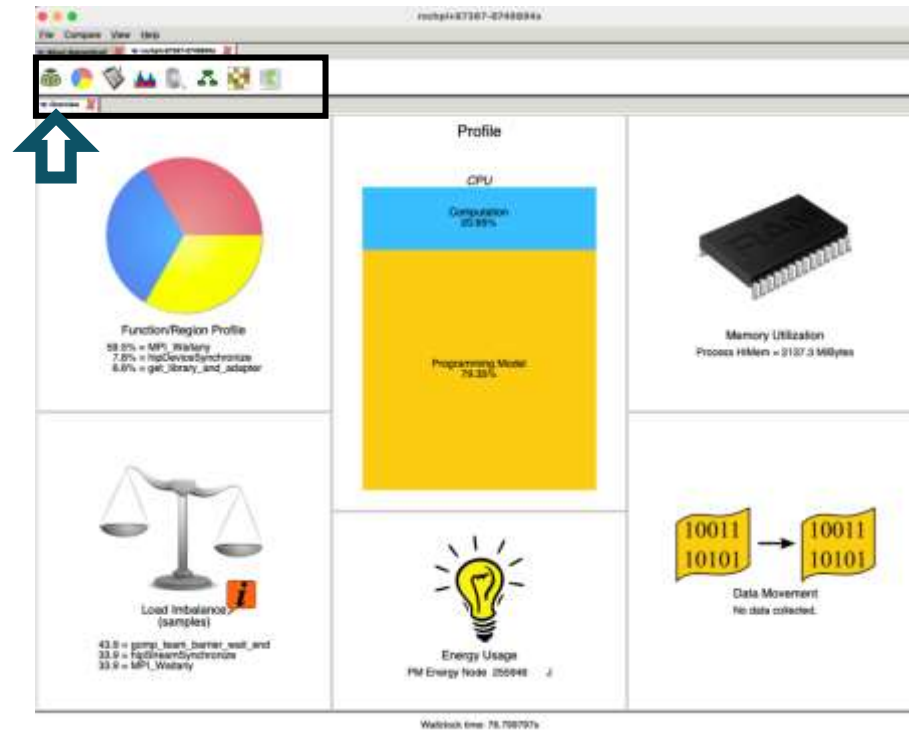
- Open a terminal and run
  - `module load perftools-base`
  - `app2 &`
- Now you can use Apprentice2 directly on LUMI
  - Open the PAT directory and analyse the data
- Also, you can use the app2 CLI, e.g.
  - Open a directory: `app2 <experiment_data_directory> &`
  - Compare two experiments: `app2 --compare <experiment_data_directory1> <experiment_data_directory2> &`
  - More info via `man app2`
- When you have finished, remember to cancel (and delete) the remote Desktop allocation





# Visualize profiling results

- Can use the Apprentice2 application:
  - `app2 main.x+*/ &`
- The actual information contained in the resulting data file will vary depending on your choice of **pat\_build** options and runtime environment variables
  - Need to run `pat_report` on the experiment directory (it produces the ap2 summary files)



- **Help** -> **Panel Help** opens a PDF with the documentation on Apprentice2



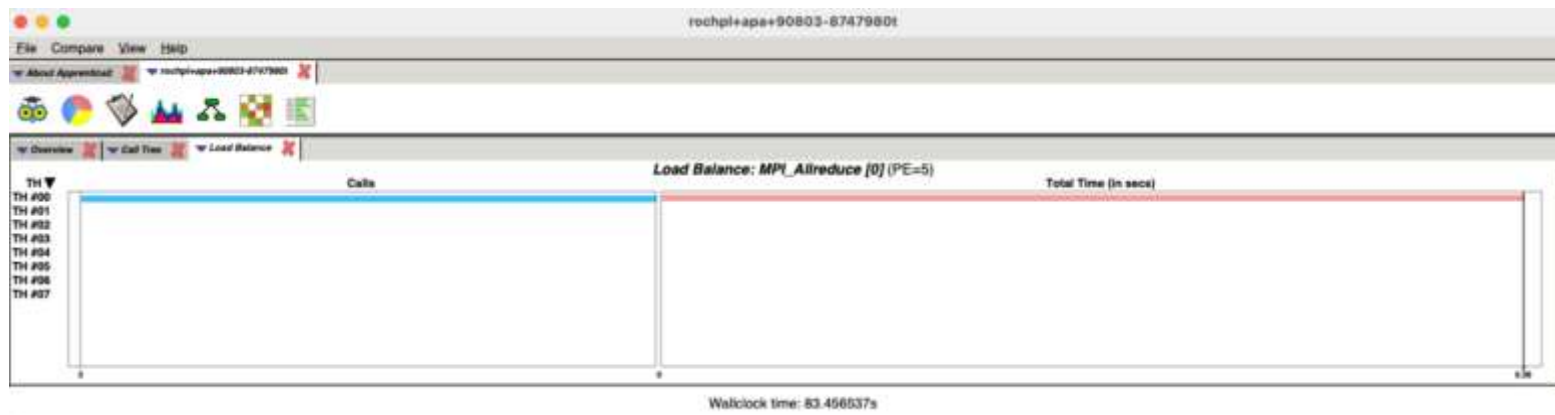


# Load Balance

- Clicking on a node of the Call tree view will show the breakdown per each PE, reporting the load balance

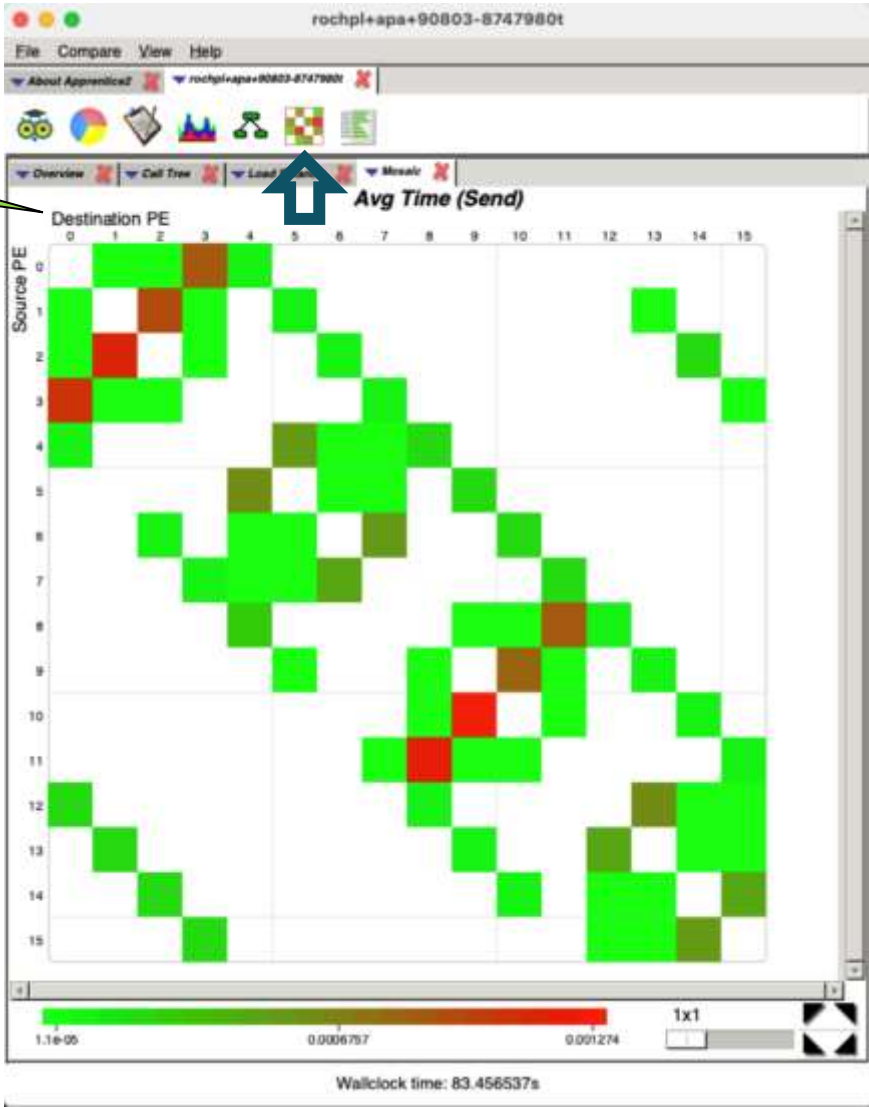


- Further clicking on a PE bar will give the breakdown per threads



# Mosaic View

Send/receive of data, useful to check communication patterns



# Experiments comparison

- Compare 2 experiments in a side-by-side display
  - Assuming they are the same e experiment and summary types
  - Useful during an application optimization phase
  - Can use CLI or the GUI menu **Compare**
- Check also the **pat\_view** utility to compare performance experiments
  - **man pat\_view**
  - Can generate pdf or csv outputs



# Time Line View

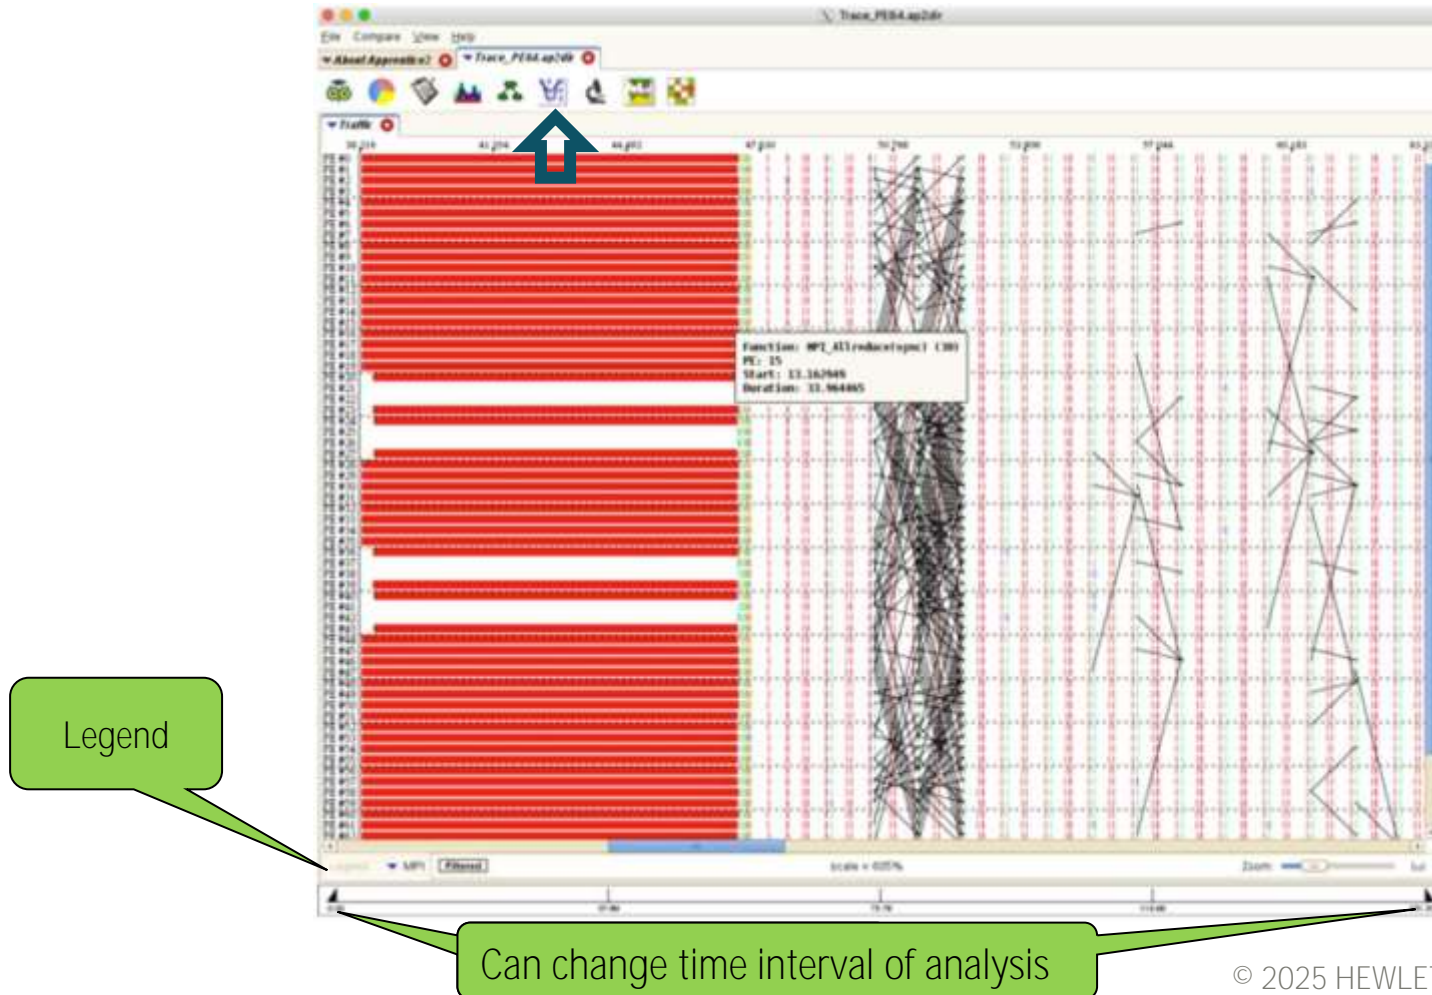
- By default, perftools aggregates data over time
- Use **PAT\_RT\_SUMMARY=0** to turn off aggregation
  - This will give you data as a function of time in Apprentice2
  - Do not run large experiments with this setting because size of data files!
- It shows:
  - Traffic time line
  - GPU time line with concurrent activity on the CPU (host) and GPU (accelerator)





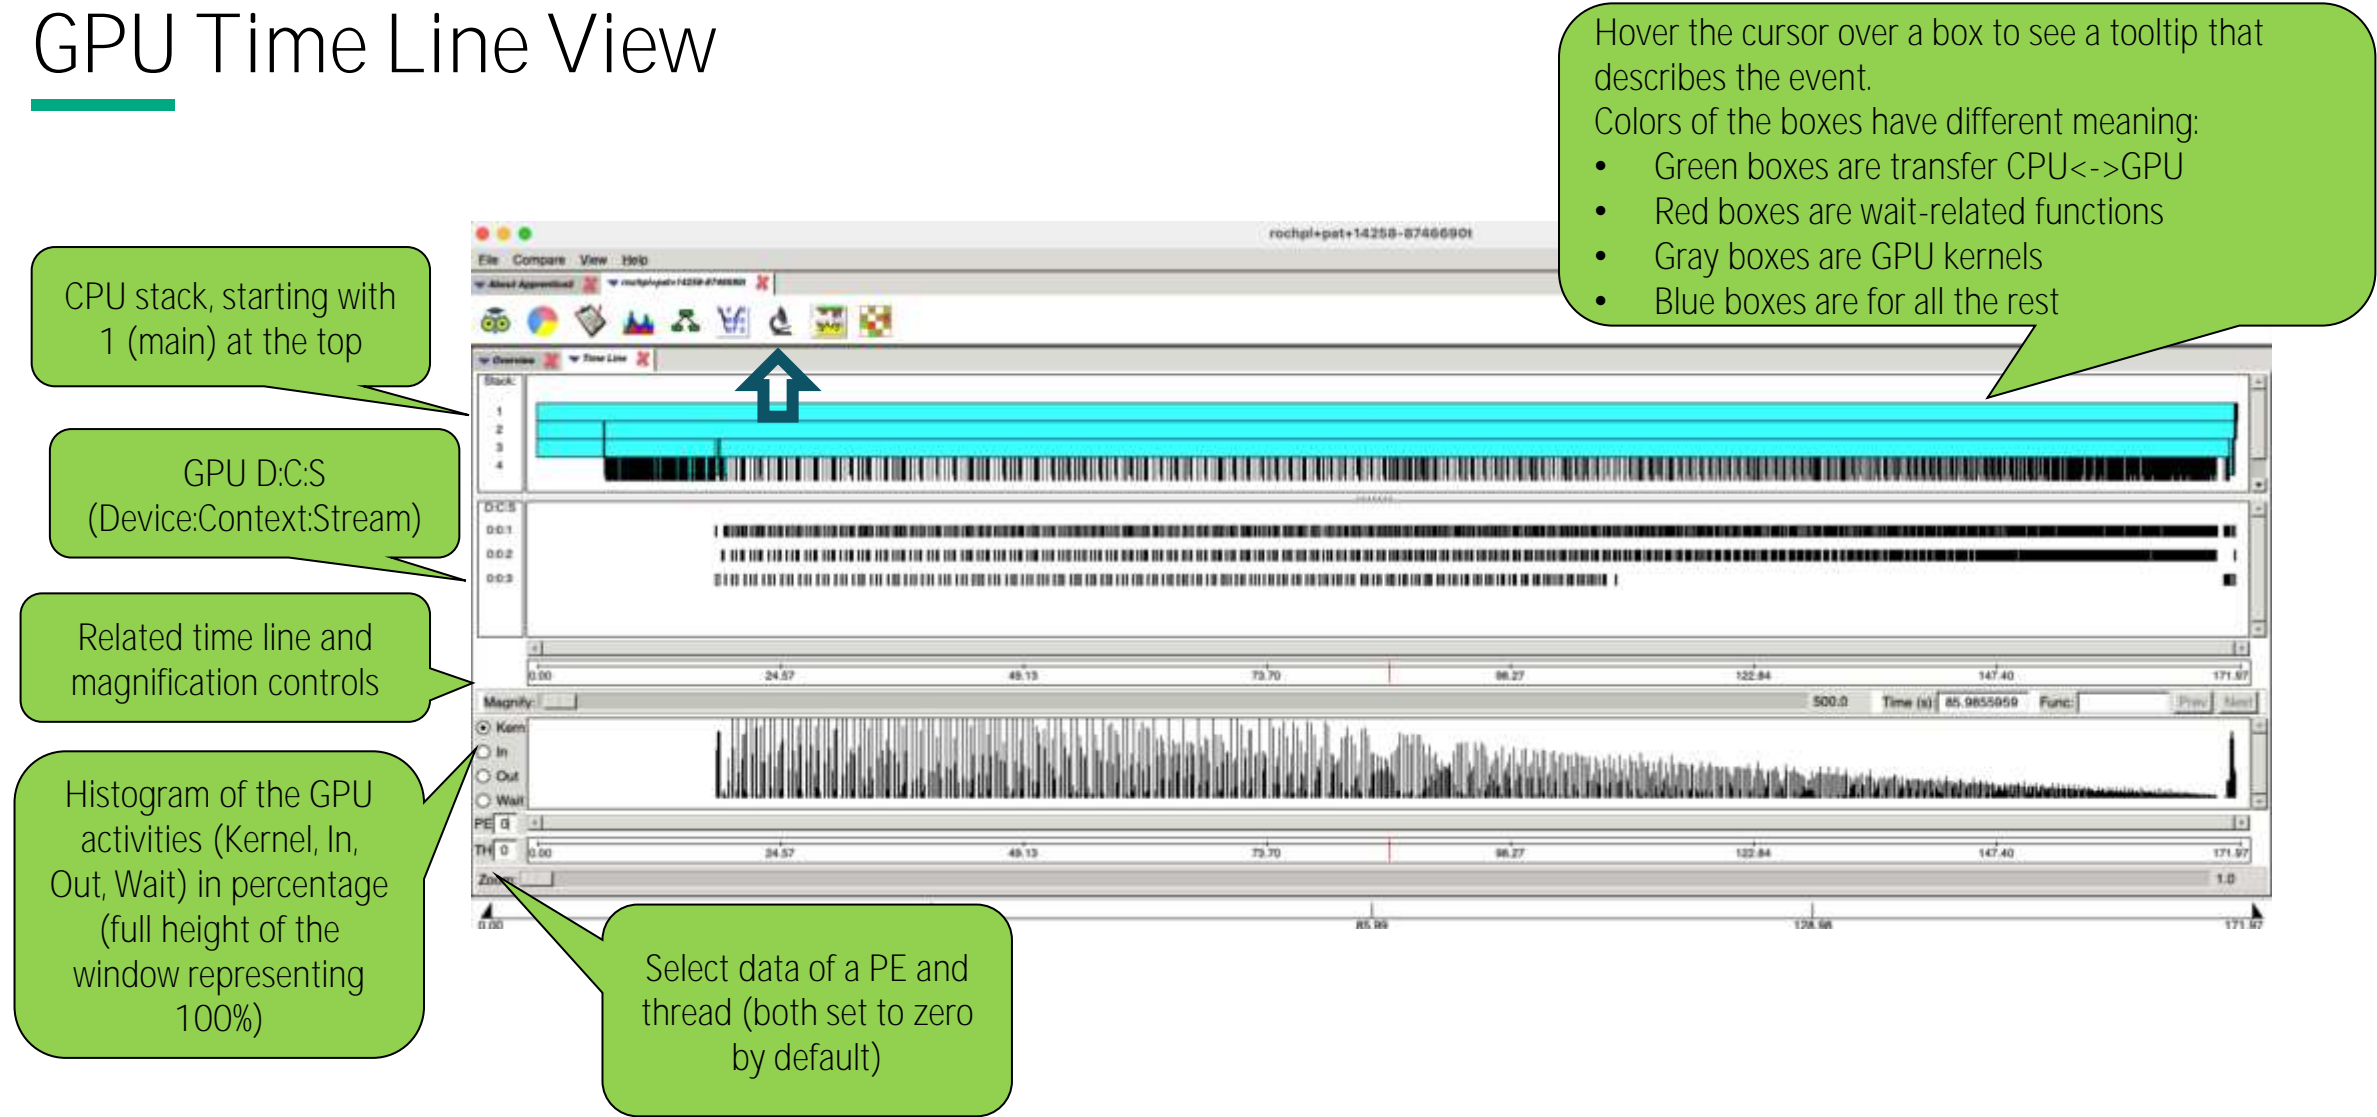
# Traffic Time Line View

- It shows internal PE-to-PE traffic and I/O activity by PE over time
  - Quite dense and typically requires zooming in or filtering to reveal meaningful data





# GPU Time Line View





# Questions?