

Hewlett Packard
Enterprise

Performance Optimization: Improving Single-core Efficiency

LUMI Advanced Workshop

March 5–7, 2025

Context of This Talk



- Scope of the presentation
 - Poor performance of an application can have many causes: placement, OS, compilers, libraries, hardware, etc.
 - We will consider possible techniques to optimize execution on a single core
 - In practice this means we should consider if we are using the processor hardware and memory hierarchy efficiently and if the compiler is compiling our application for optimum performance
- Processors have seen numerous evolutions in the last few years
 - (slight) increase in frequency, wider vector registers, execution pipelines
 - Imbalance between computational and the memory performance (memory wall)
- Important for programmers
 - Optimize performance-critical parts and be sure to use representative test-cases and problem sizes
 - Optimize data movement
 - Help the compiler generate efficient instruction sequences

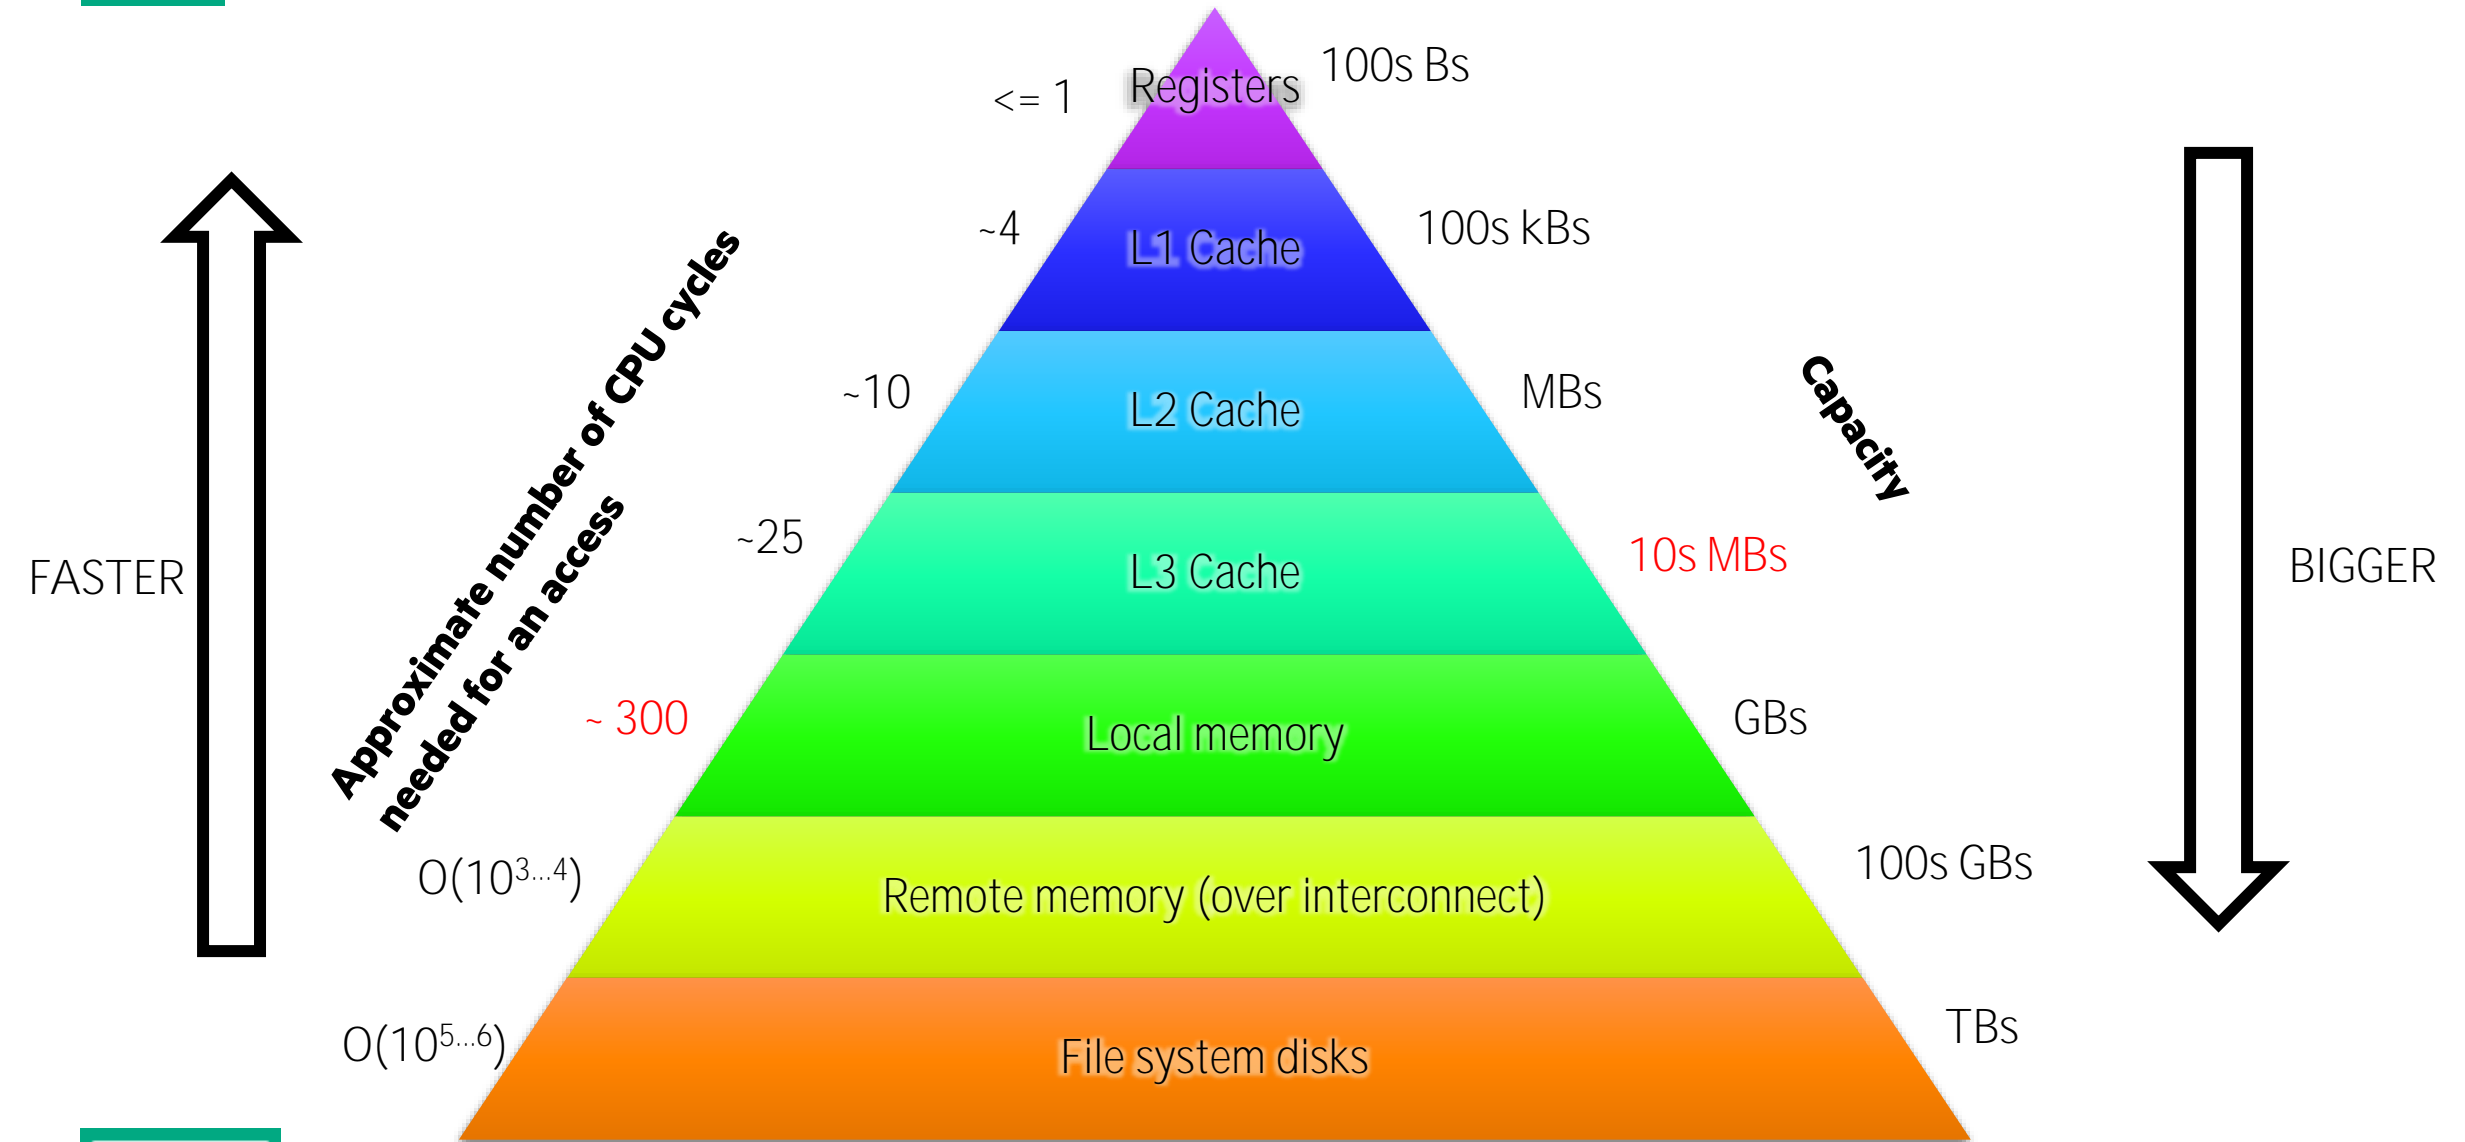


Doesn't the Compiler do Everything?

- Not yet...
 - Standard answer, unchanged for last 50 or so years
- You can make a big difference to code performance
 - Helping the compiler spot optimization opportunities
 - Using the insight of your application
 - **Removing obscure (and obsolescent) “optimizations” in older code**
 - Simple code is the best, until proven otherwise
- No fixed rules: optimize on case-by-case basis
 - But first, check what the compiler is already doing
- What we cover in this talk:
 - Loop optimization: Loop Unrolling, Strip Mining
 - Cache optimization: Cache Blocking, Data Alignment
 - Vectorization: compiler hints



Keep Your Friends Close and Data Even Closer




System Latencies

Table 2.2 Example Time Scale of System Latencies

Event	Latency	Scaled
1 CPU cycle	0.3 ns	1 s
Level 1 cache access	0.9 ns	3 s
Level 2 cache access	2.8 ns	9 s
Level 3 cache access	12.9 ns	43 s
Main memory access (DRAM, from CPU)	120 ns	6 min
Solid-state disk I/O (flash memory)	50–150 µs	2–6 days
Rotational disk I/O	1–10 ms	1–12 months
Internet: San Francisco to New York	40 ms	4 years
Internet: San Francisco to United Kingdom	81 ms	8 years
Internet: San Francisco to Australia	183 ms	19 years
TCP packet retransmit	1–3 s	105–317 years
OS virtualization system reboot	4 s	423 years
SCSI command time-out	30 s	3 millennia
Hardware (HW) virtualization system reboot	40 s	4 millennia
Physical system reboot	5 m	32 millennia

"Systems Performance: Enterprise and the Cloud" by Brendan Gregg




<https://benjdd.com/>

DATA LOCALITY




Cache Blocking

- Cache blocking is a combination of strip mining and loop interchange, designed to increase data reuse.
 - Takes advantage of temporal locality: re-reference array elements already referenced
 - Good blocking will take advantage of spatial locality: work with the cache lines!
- Many ways to block any given loop nest
 - Which loops get blocked?
 - What block size(s) to use?
- Analysis can reveal which ways are beneficial
 - How big is your cache?
 - L1 is 32 KB
 - How many cache lines can it hold?
 - each line typically 64B, so
 - How many cache lines are needed per loop iteration?
 - ...
- But trial-and-error is probably faster
 - or autotuning of the code



Cache Blocking




Cache line counter: 120

```
DO j = 1, 8
  DO i = 1, 16
    a = u(i-1,j) + u(i+1,j) &
      + u(i,j-1) + u(i,j+1) &
      - 4*u(i,j)
```

(column-major)

- Imagine a CPU architecture where:
 - each cache line holds 4 array elements
 - cache can hold 12 lines of data (48 cells)
 - Each execution of i-loop needs:
 - $3 * \text{CEILING}[(16+2)/4] = 15$ cache lines
- No cache reuse b/w j-loop iterations
 - Because 15 is greater than 12
 - iteration j loaded $u(i:i+3, j+1)$ (4 elements)
 - iteration j+1 could reuse this (for central term)
 - but it's already been evicted from the cache
- Cache misses per loopnest iteration
 - $8 * 15 = 120$

Blocking to Increase Reuse



- Block the inner loop

```
DO ib = 1, 16, 4
```

```
DO j = 1, 8
```

```
DO i = ib, ib + 4-1
```


```
  a = u(i-1,j) + u(i+1,j) &  
    + u(i,j-1) + u(i,j+1) &  
    - 4*u(i,j)
```

- Cache lines per j-loop iteration:
 - $3 * \text{CEILING}[(4+2)/4] = 6$ cache lines
 - So can hold data for 4 j-values in cache
 - because $(4+2) * \text{CEILING}[(4+2)/4] = 12$
- Cache lines per ib-loop iteration
 - $(8+2) * \text{CEILING}[(4+2)/4] = 20$
- Cache misses per loopnest iteration
 - $(16/4) * 20 = 80$
 - reduced from 120
- Better temporal locality

Cache line counter:

80

Cache Blocking



Cache line counter:

60

- Iterate over 4×4 blocks (or "tiles")
DO $jb = 1, 8, 4$
DO $ib = 1, 16, 4$
DO $j = jb, jb + 4 - 1$
DO $i = ib, ib + 4 - 1$
 $a = u(i-1, j) + u(i+1, j) +$
 $u(i, j-1) + u(i, j+1) -$
 $4 * u(i, j)$
- Cache lines per tile:
 - $(4+2) * \text{CEILING}[(4+2)/4] = 12$
 - Can reuse for (some of) next tile
- Cache lines for each jb -iteration
 - $(4+2) * \text{CEILING}[(16+2)/4] = 30$
- Cache misses per loopnest iteration
 - $(8/4) * 30 = 60$
 - reduced from 80
 - which was reduced from 120
- Better spatial locality

Cache Blocking with Cray Directives

- **Don't modify your code to do explicit blocking**
- CCE blocks well
- Get the loopmark listing
 - Identifies which loops were blocked
 - Gives the block size the compiler chose
- **...but sometimes CCE blocks better with help**

Original loopnest	Loopnest with help	Equivalent explicit code
<pre>do k = 6, nz-5 do j = 6, ny-5 do i = 6, nx-5 ! stencil enddo enddo enddo</pre>	<pre>!dir\$ blockable(j,k) !dir\$ blockingsize(16) do k = 6, nz-5 do j = 6, ny-5 do i = 6, nx-5 ! stencil enddo enddo Enddo</pre>	<pre>do kb = 6,nz-5,16 do jb = 6,ny-5,16 do k = kb,MIN(kb+16-1,nz-5) do j = jb,MIN(jb+16-1,ny-5) do i = 6, nx-5 ! stencil enddo enddo enddo enddo enddo</pre>



Further Cache Optimizations: loop fusion

- If multiple loopnests process a large array
 - First element of array will be out of cache when start second loopnest
- Improving cache use
 - Consider fusing the loopnests
 - Completely: just have one loopnest
 - Partial: have one outer loop, containing multiple inner loops

Original code	Complete fusion	Partial fusing
<pre>do j = 1, Nj do i = 1, Ni a(i,j)=b(i,j)*2 enddo enddo do j = 1, Nj do i = 1, Ni a(i,j)=a(i,j)+1 enddo enddo</pre>	<pre>do j = 1, Nj do i = 1, Ni a(i,j)=b(i,j)*2 a(i,j)=a(i,j)+1 enddo enddo</pre>	<pre>do j = 1, Nj do i = 1, Ni a(i,j)=b(i,j)*2 enddo do i = 1, Ni a(i,j)=a(i,j)+1 enddo enddo</pre>

Further Cache Optimizations: function fusion

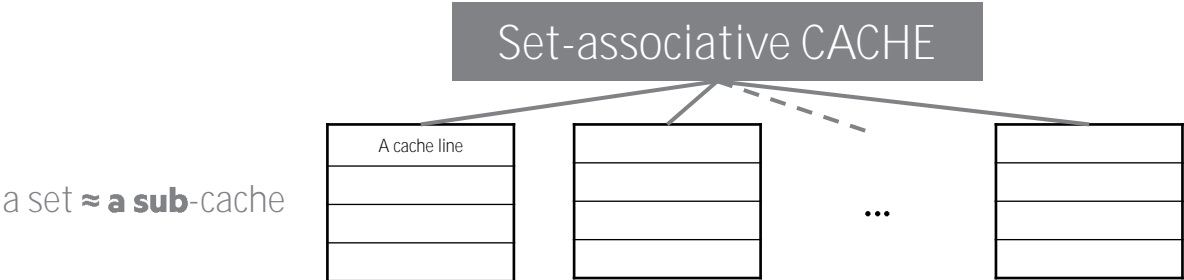
Original code	After changes
<pre>CALL sub1(a,b) CALL sub2(a) SUBROUTINE sub1(a,b) do j=1,Nj do i=1,Ni a(i,j)=b(i,j)*2 enddo enddo END SUBROUTINE sub1</pre>	<pre>do j = 1, Nj CALL sub1(a,b,j) CALL sub2(a,j) enddo SUBROUTINE sub1(a,b,j) do i=1,Ni a(i,j)=b(i,j)*2 enddo END SUBROUTINE sub1</pre>

- Perhaps cache block before fusing
 - Fuse one or more of the outer blocking loops
- If multiple subprograms process the array
 - Remove one or more outer loops (or all loops) from subprograms
 - Haul loop into parent routine, pass in index values instead
 - Might want to ensure that compiler is inlining this routine
 - This technique is useful if you want to use OpenMP/OpenACC
- Fortran note
 - With array syntax
 $a(:, :) = b(:, :) * 2$
 $a(:, :) = a(:, :) + 1$
compiler unlikely to fuse any loops



More on Increasing Reuse of Data in Cache

Reducing data cache conflict between arrays accessed in the same loop



- Adding pad arrays to align data in caches (changing memory layout)

Original code	With padding arrays
<pre>REAL(8) A(64,64), B(64,64) REAL(8) C(64,64) DO i=1,64 C(i,1)= A(i,1)+b(i,1) END DO</pre>	<pre>REAL(8) A(64,64), pad1(16), B(64,64) REAL(8) pad2(16), C(64,64) DO i=1,64 C(i,1)= A(i,1)+b(i,1) END DO</pre>

- Usually compiler does automatic padding, but you may be able to get better results adding your own pad arrays.



Cache Stats

CrayPAT profiling with `export PAT_RT_HWPC=2` (L1 and L2 metrics)

Time%		0.2%	
Time		0.000003	
Calls		1	
PAPI_L1_DCA	455.433M/sec	1367	ops
DC_L2_REFILL_MOESI	49.641M/sec	149	ops
DC_SYS_REFILL_MOESI	0.666M/sec	2	ops
BU_L2_REQ_DC	74.628M/sec	224	req
User time	0.000 secs	7804	cycles
Utilization rate		97.9%	
L1 Data cache misses	50.308M/sec	151	misses
LD & ST per D1 miss		9.05	ops/miss
D1 cache hit ratio		89.0%	
LD & ST per D2 miss		683.50	ops/miss
D2 cache hit ratio		99.1%	
L2 cache hit ratio		98.7%	
Memory to D1 refill	0.666M/sec	2	lines
Memory to D1 bandwidth	40.669MB/sec	128	bytes
L2 to Dcache bandwidth	3029.859MB/sec	9536	bytes



Better Cache Locality

Time%		0.1%	
Time		0.000002	
Calls		1	
PAPI_L1_DCA	689.986M/sec	1333 ops	
DC_L2_REFILL_MOESI	33.645M/sec	65 ops	
DC_SYS_REFILL_MOESI		0 ops	
BU_L2_REQ_DC	34.163M/sec	66 req	
User time	0.000 secs	5023 cycles	
Utilization rate		95.1%	
L1 Data cache misses	33.645M/sec	65 misses	
LD & ST per D1 miss		20.51 ops/miss	
D1 cache hit ratio		95.1%	
LD & ST per D2 miss		1333.00 ops/miss	
D2 cache hit ratio		100.0%	
L2 cache hit ratio		100.0%	
Memory to D1 refill		0 lines	
Memory to D1 bandwidth		0 bytes	
L2 to Dcache bandwidth	2053.542MB/sec	4160 bytes	





VECTORIZATION (DATA LEVEL PARALLELISM)

Single Instruction, Multiple Data (SIMD):
CPU instructions (SSE, AVX)... but also GPUs.



When does the Compiler Vectorize

- What can be vectorized
 - Mostly loops (+ SLP vectorization a.k.a. superword-level parallelism)
- Usually only one loop is vectorizable in loopnest
 - And most compilers only consider inner loop
- Optimizing code to allow compilers to use vector instructions
 - Relies on code being vectorizable
 - Or in a form that the compiler can convert to be vectorizable
 - Some compilers are better at this than others
- Check the compiler output listing and/or assembler listing
 - Look for packed SSE/AVX instructions



LOOP UNROLLING



Loop Unrolling (Theory)

Original code	After partial unrolling
<pre>do i=1,N a(i)=a(i) + b(i) enddo</pre>	<pre>do i=1,N,4 a(i) =a(i) + b(i) a(i+1)=a(i+1) + b(i+1) a(i+2)=a(i+2) + b(i+2) a(i+3)=a(i+3) + b(i+3) enddo <cleanup if N%4!=0></pre>

- Increases the work per loop iteration
 - higher computational intensity
 - more floating-point operations per memory operation (load or store)
 - more computation per iteration
 - can pipeline better in CPU
 - more opportunities for vectorization
- Combination of loop blocking and unwinding
- may completely unwind the loop



Loop Unrolling (Reality)

- Most compilers will unroll loops automatically
 - But probably will concentrate on inner loops
- When might we help?
 - When the compiler didn't unroll and should have done so
 - When the compiler doesn't know about tripcounts (this loop is small)
 - When we have a small outer loop
 - Maybe move it to be the innermost loop and completely unroll
- Avoid unrolling loops manually
 - Reduces portability
 - Compiler will likely reroll the loop and unroll it again
- Instead help compiler to do it using directives, e.g.:
 - CCE specify loop count = n if it can be estimated
!DIR\$ loop_info n
 - CCE: Force unrolling loop, optional **i** times.
!DIR\$ unroll (i)
#pragma unroll i



COMPILER AND VECTORIZATION



Helping Vectorization

- Is there a good reason for this?
 - There is an overhead in setting up vectorization; maybe it's not worth it
 - Could you unroll inner (or outer) loop to provide more work?
- Does the loop have dependencies?
 - information carried between iterations
 - e.g., counter: **total = total + a(i)**
 - No:
 - Tell the compiler that it is safe to vectorize
 - **!dir\$ IVDEP** directive above loop (CCE, but works with most compilers)
 - C99: restrict keyword
 - Yes:
 - Rewrite code to use algorithm without dependencies, e.g.
 - promote loop scalars to vectors (single dimension array)
 - use calculated values (based on loop index) rather than iterated counters, e.g.
 - Replace: **count = count + 2; a(count) = ...**
 - By: **a(2*i) = ...**
 - move **if** statements outside the inner loop
 - may need temporary vectors to do this (otherwise use [masking operations](#))
 - If you need to do too much extra work to vectorize, may not be worth it.

Code Restructuring Example

```
65.      1      PF = 0.0
66.  + 1 2-----<      DO 44030 I = 2, N
67.      1 2      AV  = B(I) * RV
68.      1 2      PB  = PF
69.      1 2      PF  = C(I)
70.      1 2      IF ((D(I) + D(I+1)) .LT. 0.) PF = -C(I+1)
71.      1 2      AA  = E(I) - E(I-1) + F(I) - F(I-1)
72.      1 2      1    + G(I) + G(I-1) - H(I) - H(I-1)
73.      1 2      BB  = R(I) + S(I-1) + T(I) + T(I-1)
74.      1 2      1    - U(I) - U(I-1) + V(I) + V(I-1)
75.      1 2      2    - W(I) + W(I-1) - X(I) + X(I-1)
76.      1 2      A(I) = AV * (AA + BB + PF - PB + Y(I) - Z(I)) + A(I)
77.      1 2-----> 44030 CONTINUE
```

Loop-carried dependency


ftn-6254 ftn: VECTOR LP44030, File = lp44030.f, Line = 66

A loop starting at line 66 was not vectorized because a recurrence was found on "pf" at line 69.



Code Restructuring Example

```
101.      1      VPF(1) = 0.0
102.      1 Vr2---<      DO 44031 I = 2, N
103.      1 Vr2          AV      = B(I) * RV
104.      1 Vr2          VPF(I) = C(I)
105.      1 Vr2          IF ((D(I) + D(I+1)) .LT. 0.) VPF(I) = -C(I+1)
106.      1 Vr2          AA      = E(I) - E(I-1) + F(I) - F(I-1)
107.      1 Vr2          1      + G(I) + G(I-1) - H(I) - H(I-1)
108.      1 Vr2          BB      = R(I) + S(I-1) + T(I) + T(I-1)
109.      1 Vr2          1      - U(I) - U(I-1) + V(I) + V(I-1)
110.      1 Vr2          2      - W(I) + W(I-1) - X(I) + X(I-1)
111.      1 Vr2          A(I) = AV * (AA + BB + VPF(I) - VPF(I-1) + Y(I) - Z(I)) + A(I)
112.      1 Vr2---> 44031 CONTINUE
```



ftn-6005 ftn: SCALAR LP44030, File = lp44030.f, Line = 102

A loop starting at line 102 was unrolled 2 times.

ftn-6204 ftn: VECTOR LP44030, File = lp44030.f, Line = 102

A loop starting at line 102 was vectorized.

Quick check:
2.2x – 3x faster



When does the Cray Compiler Vectorize?

- The Cray compiler vectorizes loops
 - Constant strides are best, indirect addressing is bad
 - Scatter/gather operations (not efficiently implemented in AVX2)
 - Can vectorize across inlined functions
 - Needs to know loop tripcount (but only at runtime)
 - i.e., Do while -style loops will not vectorize
 - No recursion allowed
 - if you have this, consider rewriting the loop
 - If you can't vectorize the entire loop, consider splitting it
 - Think about cache line : you may need to use new structures (structure of arrays, avoid array of structures)
- Check the compiler output to see what it did
 - CCE Fortran: **-hlist=a -fsave-loopmark**
 - CCE C/C++: **-fsave-loopmark**
- Clues from CrayPAT's HWPC measurements
 - **export PAT_RT_HWPC=13** or **14** # Floating point operations SP,DP
 - Complicated, but look for ratio of operations/instructions > 1
 - expect 8 for pure AVX2 with double precision floats

CONCLUSION

The background of the slide is a dark, abstract image featuring a complex network of glowing nodes and lines, resembling a data visualization or a star map. The nodes are in various colors, including blue, green, yellow, and purple. Several empty rectangular boxes are scattered across the background, some of which are highlighted with a green border. A small green horizontal line is positioned below the word 'CONCLUSION'.

CCE Fortran Directives

The Cray compiler supports a full and growing set of directives and pragmas

Fortran Cheat Sheet:

- `!dir$ concurrent`
- `!dir$ ivdep`
- `!dir$ interchange`
- `!dir$ unroll`
- `!dir$ loop_info [max_trips] [cache_na] ...` Many more
- `!dir$ blockable`

More info:

- `man directives` (only Fortran part)
- `man loop_info`

OpenMP is adding similar controls:

For example, OpenMP 5.1
[Loop Transformations](#) (tile/unroll)



Concluding Remarks

- Understand what the compiler does
 - Use hardware counters and profiling
 - Ask for the compiler feedback
- And remember what we said at the start
 - You should only attempt to optimize performance-critical parts and be sure to use representative test-cases and problem sizes
- Help the compiler to understand your code
 - Simple constructs, most basic optimizations usually better to leave for the compiler
 - Avoid branches in loops
 - Avoid indirect addressing
 - Avoid non-constant striding in loops





Questions?